

Cloud Databases: A Systematic Meta-Survey of Current Challenges, Emerging Paradigms, and Future Research Directions

Abdulhakeem Ibrahim^{a,*}

^a*School of Computer Science and Informatics, De Montfort University, Leicester, United Kingdom*

ARTICLE INFO

Keywords:

cloud databases
systematic meta-survey
cloud-native databases
database-as-a-service
data lakehouse
AI4DB
survey of surveys

ABSTRACT

Cloud databases have become the backbone of modern data-driven applications, enabling organizations to store, process, and analyze data at unprecedented scale and elasticity. Over more than two decades (1999–2026), 85 survey papers satisfying systematic inclusion and quality criteria have examined various facets of cloud database technology—spanning cloud-native architectures, hybrid transactional and/or analytical processing (HTAP), NoSQL/NewSQL systems, vector and graph databases, AI-for-database optimization (AI4DB), data lakehouses, time-series databases, cloud database migration, and edge–cloud data management. However, the challenges, open problems, and research directions identified in these surveys remain fragmented across disparate communities and publication venues. This paper presents the first *systematic meta-survey* of cloud database research: a survey of surveys that synthesizes findings from carefully selected survey papers using a rigorous methodology grounded in the Kitchenham–Charters systematic literature review protocol. The scattered landscape of cloud database challenges is organized into a structured, dual-theme taxonomy: (1) *eighteen cross-cutting challenges* that transcend individual system categories, including scalability, consistency trade-offs, security, cost optimization, AI–database convergence, and sustainability; and (2) *five lifecycle-phase challenges* spanning design, deployment, operation, optimization, and evolution of cloud database systems. The convergence of cloud databases with large language models (LLMs) is further examined as an illustrative case study, and a consolidated set of open research directions is distilled from the cross-survey synthesis. The result is a single, thematically organized map of two decades of cloud database research—one that newcomers can use to orient themselves in the literature and that system builders and researchers can use to locate the open problems where new work is most needed.

1. Introduction

The cloud database market has experienced explosive growth over the past decade, and industry analysts consistently project continued double-digit growth in cloud database spending [1]. This growth reflects a fundamental shift in how organizations manage data: from on-premises, monolithic database servers to elastic, globally distributed, cloud-native data platforms that can scale storage and compute independently on demand [2, 3]. Major cloud providers—Amazon Web Services (AWS), Microsoft Azure, Google Cloud Platform (GCP), and Alibaba Cloud—now offer dozens of specialized database services, from relational engines like Amazon Aurora [2] and Azure SQL [4] to purpose-built systems for documents, graphs, time series, key-value pairs, and, most recently, vector embeddings for AI workloads [5].

This rapid evolution has attracted significant research attention. A scripted search across four scholarly indexing services (OpenAlex, arXiv, Crossref, and Semantic Scholar) covering January 1999 to July 2026 identifies **85 survey papers** meeting the inclusion and quality criteria, each examining a particular facet of cloud database technology. These surveys span a wide spectrum of topics: cloud-native database architectures [6], hybrid transactional/analytical processing (HTAP) [7], graph databases [8], vector database management systems [5], AI-driven database optimization (AI4DB) [9], data lakehouses [10], NoSQL and NewSQL systems [11, 12], edge and fog database systems [13], and many more.

Despite this wealth of individual surveys, the field lacks a *meta-level synthesis*—a comprehensive, systematic effort to consolidate the challenges, open problems, and research directions scattered across these surveys into a unified framework. Each survey naturally focuses on its own domain, leading to a fragmented understanding of the cloud

*Corresponding author.

✉ hakeem.ibrahim@dmu.ac.uk (A. Ibrahim)
ORCID(S): 0000-0002-1104-5314 (A. Ibrahim)

Table 1

Research questions guiding the meta-survey.

ID	Research Question
RQ1	What is the landscape of cloud database surveys published between 1999 and 2026, in terms of topic coverage, methodology, and publication venue?
RQ2	What are the major cross-cutting challenges identified across cloud database surveys, and how do they relate to one another?
RQ3	How do cloud database challenges map to the lifecycle phases of designing, deploying, operating, optimizing, and evolving database systems?
RQ4	What is the current state and future trajectory of the convergence between cloud databases and AI/ML technologies, particularly large language models?
RQ5	What are the most pressing open research directions, and which of them recur most widely across the surveyed literature?

database landscape as a whole. Cross-cutting challenges such as cost optimization, security, consistency trade-offs, and AI–database convergence are discussed in many surveys but from different angles, without a common organizing framework.

1.1. Contributions and novelty

This paper addresses the gap by presenting the **first systematic meta-survey of cloud database research**—a survey of surveys. A rigorous systematic literature review (SLR) protocol based on the Kitchenham–Charters guidelines [14] is applied to identify, evaluate, and synthesize cloud database surveys. The contributions of this work are as follows:

- C1 Systematic identification and classification.** Four scholarly APIs are searched with a fully scripted, reproducible protocol over the window January 1999 to July 2026, yielding 85 quality-assessed survey papers on cloud databases, classified by topic, scope, methodology, and venue.
- C2 Dual-theme challenge taxonomy.** The scattered challenges are organized into a structured taxonomy with two complementary themes:
 - *Theme 1: Eighteen cross-cutting challenges* that recur across multiple survey domains, including scalability, consistency, security, cost modeling, AI–database convergence, and sustainability.
 - *Theme 2: Five lifecycle-phase challenges* that map to the stages of designing, deploying, operating, optimizing, and evolving cloud database systems.
- C3 LLM–database convergence analysis.** The rapidly emerging intersection of large language models (LLMs) and cloud databases is examined as an illustrative case study, covering text-to-SQL, vector databases for retrieval-augmented generation (RAG), and LLM-based database administration.
- C4 Research agenda.** A consolidated, cross-cutting research agenda is distilled from the synthesis, grouping the open problems that recur most widely across the surveyed literature.
- C5 Knowledge engineering perspective.** Cloud database challenges are explicitly connected to data and knowledge engineering concerns—including knowledge graphs, ontologies, metadata management, and schema evolution—making this survey particularly relevant to the *Data & Knowledge Engineering* community.

1.2. Research questions

This meta-survey is guided by five research questions, summarized in Table 1 and elaborated in Section 4.

1.3. Paper organization

The remainder of this paper is organized as shown in Figure 1. Section 2 motivates the importance of cloud databases from five complementary perspectives. Section 3 establishes key definitions and distinctions—from cloud-hosted to cloud-native databases, architectural paradigms, data models, and the data platform continuum. Section 4 details the systematic review methodology, including search strategy, selection criteria, quality assessment, and results. Section 5 presents the core contribution: a comprehensive discussion of challenges and research directions organized by cross-cutting themes and lifecycle phases. Section 6 examines the convergence of cloud databases and LLMs. Section 7 concludes. Appendix A provides the complete catalog of surveyed papers.

Introduction	[Section 1]
Why Cloud Databases Matter	[Section 2]
From Cloud Databases to Cloud-Native Data Platforms	[Section 3]
Systematic Review Planning and Execution	[Section 4]
Discussion—Challenges and Research Directions	[Section 5]
Theme 1: Eighteen cross-cutting challenges	
Theme 2: Five lifecycle-phase challenges	
Cloud Databases and LLMs—A Convergence Case Study	[Section 6]
Conclusions	[Section 7]

Figure 1: The organization of this meta-survey paper.

2. Why Cloud Databases Matter—A Multi-Perspective Motivation

Understanding the significance of cloud databases requires examining the topic from multiple viewpoints. In this section, the need for a meta-survey is motivated by discussing five complementary perspectives (Figure 2), each revealing why the scattered state of survey knowledge is problematic and why a unified synthesis is timely.




Figure 2: The five complementary perspectives motivating a systematic meta-survey of cloud databases.

2.1. Industry perspective

Analyst assessments of the database management system (DBMS) market agree on two structural trends: cloud deployments now account for well over half of total DBMS revenue, and their share continues to grow at the expense of on-premises deployments [1]. Gartner’s widely cited 2019 projection that three quarters of all databases would be deployed or migrated to cloud platforms by 2022 has, by most industry accounts, been broadly borne out [1]. This shift is driven by digital transformation initiatives across industries, from financial services adopting cloud databases for real-time fraud detection to healthcare organizations leveraging cloud data platforms for genomics and precision medicine.

The vendor landscape has also diversified significantly. Beyond the hyperscale cloud providers (AWS, Azure, GCP), specialized cloud database vendors such as Snowflake [3], Databricks [15], CockroachDB [16], PlanetScale, and Neon have gained substantial market share. China’s cloud database ecosystem—including Alibaba PolarDB [17], TiDB [18], and OceanBase [19]—has matured rapidly. This fragmentation makes it difficult for practitioners to navigate the landscape without a systematic overview.

2.2. Data engineering perspective

Modern data architectures have evolved from simple transactional databases to complex, multi-layered data ecosystems. The emergence of the *data lakehouse* paradigm [15] exemplifies this evolution, combining the flexibility of data lakes with the governance and performance of data warehouses. Systems like Databricks Delta Lake, Apache Iceberg, and Apache Hudi have created a new “open table format” layer that bridges batch and streaming analytics [10].

Simultaneously, the rise of *HTAP* (Hybrid Transactional/Analytical Processing) systems [7] challenges the traditional separation between OLTP and OLAP workloads. Systems such as TiDB [18], SAP HANA [20], and Google AlloyDB promise unified engines that handle both transaction processing and analytical queries. Understanding the trade-offs, maturity, and open challenges of these paradigms requires synthesizing insights across multiple survey papers.

2.3. Knowledge engineering perspective

Cloud databases play a crucial role in knowledge engineering, making this topic directly relevant to the *Data & Knowledge Engineering* (DKE) community. Several connections are worth highlighting:

- **Knowledge graphs in the cloud.** Large-scale knowledge graphs (KGs) such as Google Knowledge Graph, Wikidata, and enterprise KGs are increasingly stored and queried using cloud-hosted graph databases (e.g.,

Amazon Neptune, Azure Cosmos DB with Gremlin API, Neo4j Aura) [8, 21, 22]. The scalability, consistency, and query optimization challenges of cloud graph databases directly impact KG applications.

- **Ontology-mediated data access.** Cloud environments with heterogeneous data sources benefit from ontology-based data integration (OBDI) techniques [23], where ontologies provide a unified conceptual layer over distributed cloud databases.
- **Metadata and schema management.** Cloud data catalogs and metadata management systems (e.g., AWS Glue Data Catalog, Apache Atlas) implement knowledge engineering principles for organizing, discovering, and governing data assets [24].
- **AI-knowledge engineering convergence.** The integration of machine learning with database systems (AI4DB) [9] draws on knowledge engineering concepts such as learned models for query optimization, workload characterization, and automated tuning.

2.4. Economic and cost perspective

Cloud databases introduce novel economic dynamics compared to on-premises deployments. The shift from capital expenditure (CapEx) to operational expenditure (OpEx) models, combined with pay-per-use pricing and complex multi-dimensional billing (compute, storage, I/O, network egress, backups), creates significant challenges in cost prediction and optimization [25, 26].

Research has shown that cloud database costs can vary by orders of magnitude depending on configuration choices, query patterns, and pricing model selection [25]. The total cost of ownership (TCO) analysis becomes particularly complex for multi-cloud or hybrid cloud deployments. Several surveys touch on cost-related challenges [6, 27], but a consolidated analysis is lacking.

2.5. Regulatory and compliance perspective

Data sovereignty regulations such as the European Union’s General Data Protection Regulation (GDPR) [28], China’s Data Security Law, and sector-specific requirements (e.g., HIPAA for healthcare, PCI-DSS for payment data) impose strict constraints on how and where cloud databases store and process data [29].

Multi-region deployment, data residency requirements, cross-border data transfer restrictions, and audit trail obligations all impact cloud database architecture and operations. The compliance landscape interacts with technical challenges such as geo-distributed consistency, encryption, access control, and data lineage—topics discussed in various surveys but rarely addressed in an integrated manner.

3. From Cloud Databases to Cloud-Native Data Platforms

Before analyzing the surveyed literature, key concepts, definitions, and distinctions are established that form the vocabulary of cloud database research.

3.1. Cloud-hosted vs. cloud-native vs. cloud-born databases

The term “cloud database” encompasses a spectrum of architectures with fundamentally different design philosophies. Three categories are distinguished, compared in Table 2:

Cloud-hosted databases are traditional database engines (e.g., PostgreSQL, MySQL, Oracle) deployed on cloud virtual machines. They benefit from cloud infrastructure (networking, storage volumes, automated backups) but do not exploit cloud-native primitives such as elastic compute pools or disaggregated storage [6].

Cloud-native databases are existing or new engines that have been substantially re-architected to leverage cloud services. Amazon Aurora [2] exemplifies this approach: it re-implements the storage layer of MySQL/PostgreSQL to use a distributed, log-structured storage service while keeping the SQL engine largely intact. Microsoft’s Socrates [4] takes SQL Server and disaggregates its components across cloud microservices.

Cloud-born databases are designed entirely from the ground up for cloud deployment, with no legacy on-premises heritage. Snowflake [3] pioneered this model for data warehousing, while CockroachDB [16] and Neon represent cloud-born approaches for transactional and PostgreSQL-compatible workloads, respectively.

Table 2

Comparison of cloud database deployment models.

Characteristic	Cloud-Hosted	Cloud-Native	Cloud-Born
Definition	Traditional DBMS deployed on cloud VMs	DBMS redesigned to exploit cloud infrastructure	DBMS designed from scratch for the cloud
Architecture	Monolithic, tightly coupled	Disaggregated compute/storage	Fully disaggregated, microservice-based
Elasticity	Manual scaling (vertical)	Automatic scaling (horizontal)	Instant, fine-grained scaling
Examples	MySQL on EC2, Oracle on Azure VM	Amazon Aurora [2], Azure SQL [4]	Snowflake [3], Neon, CockroachDB [16]
State sharing	Shared-nothing (typically)	Shared-storage	Disaggregated, shared cloud services
Multi-tenancy	VM-level isolation	Engine-level isolation	Native multi-tenant design

3.2. A brief history: from early outsourced databases to cloud-native platforms (1999–2015)

The systematic search window of this meta-survey opens in January 1999, yet the earliest survey satisfying all inclusion and quality criteria dates from 2009 (Table 11)—a finding that itself traces the field’s history. The idea of consuming database functionality as an outsourced service predates the cloud era: the “database-as-a-service” model was articulated in the early 2000s around outsourced, encrypted query processing, and the security and confidentiality questions it raised dominated the first wave of survey activity [30, 31]. Commercial cloud databases arrived with Amazon SimpleDB and Google Bigtable [32], whose early comparative assessments [33] and elasticity measurements [34] mark the transition from outsourced databases to genuinely cloud-resident data stores. The NoSQL wave of 2009–2013 [35, 36] multiplied data models and triggered the first generation of comparative NoSQL surveys, while practitioner-oriented overviews began mapping the emerging cloud database service landscape [37, 38, 39]. By the mid-2010s, quality-attribute-driven selection studies [40] and consolidated DBaaS security taxonomies [31] signal a maturing field—setting the stage for the cloud-native re-architecting era that the majority of the catalog documents from 2016 onward.

3.3. Database-as-a-Service (DBaaS) spectrum

Database-as-a-Service (DBaaS) represents the fully managed end of the cloud database spectrum, where the cloud provider handles provisioning, patching, scaling, backup, and high availability [41, 42, 43, 44, 37]. The DBaaS model has evolved along several dimensions:

- **Managed databases:** Provider-managed instances of open-source or commercial engines (e.g., Amazon RDS, Azure Database for PostgreSQL).
- **Serverless databases:** Fully elastic databases that scale to zero when idle and automatically provision resources based on demand (e.g., Amazon Aurora Serverless, Azure Cosmos DB serverless, Neon) [27].
- **Autonomous databases:** Self-driving databases that use ML to automate tuning, indexing, partitioning, and anomaly detection (e.g., Oracle Autonomous Database) [45].

3.4. Architectural paradigms


Cloud database architectures have evolved through three major paradigms, illustrated in Figure 3:

Shared-nothing architectures partition data across independent nodes, each with its own CPU, memory, and disk. Traditional distributed databases (e.g., early Spanner [46], Cassandra) follow this model. While horizontally scalable, resharding and rebalancing are expensive.


Shared-disk (cloud storage) architectures decouple compute from storage, allowing compute nodes to share a common cloud storage layer. Amazon Aurora [2] and Snowflake [3] exemplify this paradigm, enabling independent scaling of compute and storage.

Fully disaggregated architectures further decompose the database engine into independently scalable microservices—query processing, transaction management, logging, metadata management, and storage [4]. Microsoft Socrates and emerging systems like PolarDB [17] pursue this vision.

Shared-Nothing



Shared-Disk (Cloud Storage)



Fully Disaggregated




Figure 3: Evolution of cloud database architectures: from shared-nothing clusters to shared-disk (cloud storage) to fully disaggregated microservice architectures.

3.5. Data models in the cloud

Cloud platforms support a diverse ecosystem of data models, each optimized for different workload characteristics [47, 48, 36, 49, 50, 51]; early comparisons of the first-generation cloud stores such as SimpleDB and Bigtable already highlighted this model diversity [33]. Table 3 compares the major data models available in cloud database services.

The trend toward *multi-model databases* [52] reflects the desire to avoid managing separate database systems for different data models. Azure Cosmos DB, for instance, supports document, graph, key-value, column-family, and table APIs through a single backend. However, multi-model integration introduces challenges in query optimization, consistency semantics, and performance trade-offs [52].

3.6. The data platform continuum: warehouses, lakes, and lakehouses

The cloud data management landscape has converged around three complementary paradigms, illustrated in Figure 4:

Table 3

Cloud database data model landscape.

Data Model	Strengths	Cloud Examples	Typical Use Cases
Relational	ACID, SQL, mature ecosystem	Aurora, Cloud SQL, Azure SQL	OLTP, ERP, financial systems
Document	Schema flexibility, nested data	MongoDB Atlas, Cosmos DB, Firestore	Content management, catalogs, user profiles
Key-Value	Ultra-low latency, high throughput	DynamoDB, Redis Cloud, Memorystore	Caching, session management, real-time
Wide-Column	High write throughput, time-series	Bigtable, Cassandra (Astra), ScyllaDB	IoT, telemetry, event logging
Graph	Relationship traversal, pattern matching	Neptune, Neo4j Aura, Cosmos DB (Gremlin)	Social networks, fraud detection, KGs
Vector	Similarity search, ANN indexing	Pinecone, Weaviate, Milvus, pgvector	RAG, recommendation, image search
Time-Series	Temporal queries, downsampling	TimescaleDB, InfluxDB Cloud, Timestream	Monitoring, financial ticks, sensor data
Multi-Model	Multiple models in one engine	Cosmos DB, ArangoDB, SurrealDB	Polyglot persistence, unified access

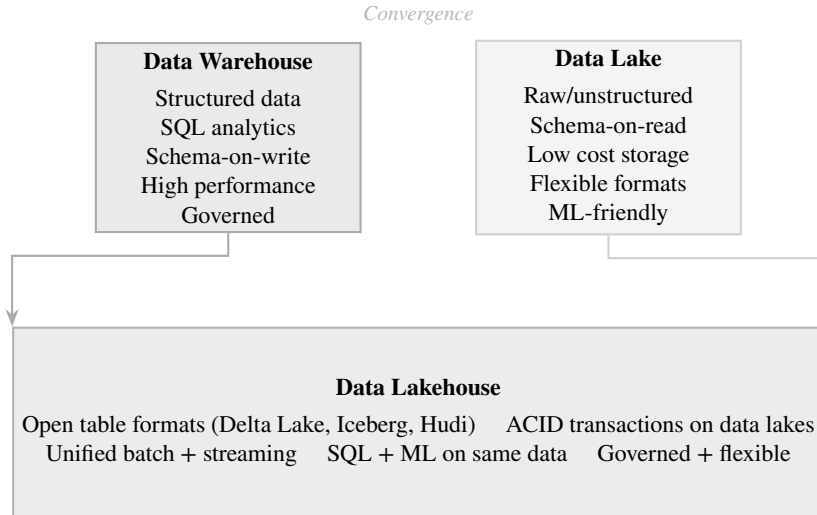


Figure 4: The data platform continuum: data warehouses and data lakes converge into the data lakehouse paradigm.

Data warehouses (e.g., Snowflake [3], BigQuery [53], Redshift) provide high-performance SQL analytics over structured, curated data with strong governance and schema enforcement.

Data lakes (e.g., Amazon S3 + Athena, Azure Data Lake Storage) offer cost-effective storage for raw, semi-structured, and unstructured data, enabling schema-on-read flexibility for data science and ML workloads.

Data lakehouses [15] combine the best of both paradigms by adding ACID transactions, schema enforcement, and data governance to open data lake storage through open table formats such as Delta Lake, Apache Iceberg, and Apache Hudi [10, 54].

Table 4

Scholarly APIs searched (5 July 2026) and raw record yields. The OpenAlex sweep issued one query per cloud-database term (18 queries), each conjoined with the survey-indicator block of the canonical search string below; the arXiv sweep combined the cs.DB category with title-level survey indicators.

API	Query scope	Records
OpenAlex	Title + abstract; 18 term-queries × survey-indicator block; 1999–2026; article/review/preprint; English	3,684
arXiv API	cat:cs.DB × title survey-indicators	356
Crossref (suppl.)	Bibliographic relevance queries; journal articles from 1999	600
Semantic Scholar (suppl.)	Bulk keyword search; 1999–2026	3,035
Total raw records		7,675
After deduplication (DOI, arXiv id, normalized-title matching)		6,613

4. Systematic Review Planning and Execution

The systematic literature review (SLR) guidelines established by Kitchenham and Charters [14], which have been widely adopted in software engineering and computer science [55, 56, 57, 58], are followed. This protocol prescribes a three-phase process: (1) *planning*, in which research questions, search strategy, and inclusion/exclusion criteria are defined; (2) *conducting*, in which studies are systematically identified, selected, and their data extracted; and (3) *reporting*, in which the synthesized findings are documented and validated through quality assessment. The selection process is additionally reported using the PRISMA 2020 statement [59], the de facto standard for transparent reporting of systematic reviews.

4.1. Search strategy

The search was executed as a fully scripted, reproducible sweep over four scholarly indexing services via their public APIs on 5 July 2026, as detailed in Table 4. OpenAlex and the arXiv API served as the primary sources; Crossref and Semantic Scholar were queried as supplementary sources to catch records missed by the primary indexes. The complete search scripts, the per-query result log, and every screening decision are archived alongside the manuscript (see the Data availability statement), so the identification stage can be re-executed verbatim.

The canonical search string combines cloud database terminology with survey-indicating terms; it was applied natively by the OpenAlex queries and re-applied uniformly to all records (including the relevance-ranked supplementary sources) as an automated eligibility marking step:

```
("cloud database" OR "cloud-native database" OR "DBaaS" OR "database-as-a-service" OR
"serverless database" OR "HTAP" OR "NewSQL" OR "data lakehouse" OR "vector database" OR
"AI4DB" OR "learned index" OR "NoSQL" OR "distributed database" OR "multi-model database"
OR "graph database" OR "edge database" OR "fog database" OR "autonomous database")
AND
("survey" OR "review" OR "tutorial" OR "taxonomy" OR "systematic" OR "state-of-the-art" OR
"overview")
```

4.2. Study selection

The inclusion and exclusion criteria listed in Table 5 were applied in a two-phase process: (1) title and abstract screening, and (2) full-text assessment.

4.3. Quality assessment

Each candidate survey was evaluated against eight quality assessment criteria (QA1–QA8), scored on a 3-point scale (0 = No, 0.5 = Partially, 1 = Yes). Given the breadth of the expanded 1999–2026 identification pool, the inclusion threshold was set at 6 out of 8; papers scoring below this threshold were excluded (papers scoring between 4 and 6 are reported explicitly in the selection flow). QA7 was scored under an explicit venue-tier rule—1 for venues of established

Table 5

Inclusion and exclusion criteria.

ID	Criterion
<i>Inclusion Criteria</i>	
IC1	The paper is a survey, review, tutorial, systematic mapping study, or cross-system comparative experimental evaluation
IC2	The paper focuses primarily on cloud-hosted, cloud-native, or cloud-relevant database technologies
IC3	The paper was published between January 1999 and the search date (July 2026)
IC4	The paper is written in English
IC5	The paper is peer-reviewed or published in a recognized venue (including arXiv with significant citations)
<i>Exclusion Criteria</i>	
EC1	The paper is a primary study (not a survey/review)
EC2	The paper surveys cloud computing in general without database-specific focus
EC3	The paper is a short paper (<6 pages) or workshop position paper
EC4	The paper is a duplicate or substantially overlapping earlier version
EC5	The full text could not be retrieved (no open-access copy and no manually obtainable copy)
EC6	The venue exhibits predatory indicators: self-claimed impact factors, discontinued or absent reputable indexing (Scopus/DOAJ), a watchlisted publisher, or an unverifiable editorial process

Table 6

Quality assessment checklist for surveyed papers.

ID	Quality Assessment Question
QA1	Are the survey objectives clearly stated?
QA2	Is the scope and coverage of the survey well-defined?
QA3	Is the search/selection methodology described or implied?
QA4	Does the survey provide a taxonomy or classification framework?
QA5	Are the reviewed primary studies adequately described?
QA6	Does the survey identify open challenges or future directions?
QA7	Is the survey published in a recognized venue (journal, top conference, or high-citation arXiv)?
QA8	Is the survey sufficiently comprehensive (≥ 30 references)?

reputation, 0.5 for minor but verifiable outlets (university, society, or currently Scopus-/DOAJ-indexed publishers), and 0 otherwise—and, independently of the QA total, papers from venues meeting the EC6 predatory-indicator test were excluded in a dedicated venue-integrity audit reported in the selection flow.

4.4. Results: study selection and demographics

Figure 5 presents the PRISMA 2020 flow diagram of the selection process, with two identification arms: records identified through the scripted database search, and records identified through other methods (the previous version of this catalog, manually collected reports, and backward snowballing). From 7,675 raw database records, 85 surveys were ultimately included in the meta-survey.




Figure 5: PRISMA 2020 flow diagram of the study selection process [59]. All figures are scripted tallies over the archived screening-decision log. The full-text exclusion box aggregates one merged duplicate of a prior-catalog entry; of the six snowball candidates, one was retrieved and excluded below the quality threshold and five could not be retrieved. Of the 24 prior-catalog entries, two failed the raised quality threshold and two were re-identified by the database search and are counted in the databases arm, leaving 20 carried forward via the other-methods arm.

Figure 6 shows the distribution of selected surveys by publication year. Survey activity is sparse before 2016, grows steadily through the late 2010s, and intensifies markedly from 2024 onward—2024 and 2025 together account for more than a third of the corpus.




Figure 6: Distribution of the 85 selected surveys by publication year. Although the search window opens in January 1999 and extends to the search date (July 2026), the earliest survey satisfying all inclusion and quality criteria dates from 2009 and the latest from 2025.

Figure 7 shows the distribution by primary topic domain (each survey is assigned a single primary topic, so the counts sum to 85). NoSQL/NewSQL systems are the most frequently surveyed topic, reflecting the maturity and breadth of that body of literature, followed by graph databases and AI4DB/learned database components.




Figure 7: Distribution of the 85 selected surveys by primary topic domain.

Figure 8 shows the distribution by venue type. Journal publications account for the majority, followed by arXiv preprints and conference proceedings; the visible arXiv share reflects both the field's preprint culture and the open-access retrieval policy described in Section 4.5.




Figure 8: Distribution of the 85 selected surveys by venue type.

4.5. Threats to validity

As with any systematic review, this meta-survey is subject to validity threats, which are mitigated as follows.

Construct validity. The search string and the choice of indexing services may not capture every relevant survey, particularly those using non-standard terminology or absent from the OpenAlex, arXiv, Crossref, and Semantic Scholar indexes; abstract availability in these indexes is also uneven, which affects the automated eligibility marking. These threats are mitigated by (i) running fully scripted queries whose logs are archived and re-executable, (ii) taking the union of four independent APIs rather than relying on a single index, (iii) carrying forward the previous version of the catalog as a separate identification arm, and (iv) a backward-snowballing pass over the reference lists of the highest-scoring included surveys, which surfaced no additional qualifying survey—evidence that the search had reached saturation. The English-language filter may exclude relevant surveys in other languages.

Internal validity. Title/abstract screening, full-text assessment, and quality-assessment scoring were performed with the assistance of a large-language-model tool operating under fixed written rubrics (disclosed in the AI-use declaration), with every record's decision and rationale logged; a strict second re-check pass was applied to all provisional inclusions, a dedicated venue-integrity audit re-examined venue quality, and all borderline and final decisions were adjudicated by the author. Every included survey was additionally verified against a resolvable DOI or arXiv identifier and a retained, content-verified full-text copy, so that the catalog contains only confirmable publications. The absence of a second independent *human* reviewer, and the reliance on AI assistance in screening, nonetheless remain limitations.

External validity. As a survey of surveys, the findings inherit the scope and currency limitations of the primary surveys; rapidly evolving topics (e.g., LLM–database convergence) may be underrepresented relative to their current activity. The corpus is further subject to an open-access retrieval bias: every included survey must have a content-verified full text on file, and 292 otherwise-eligible reports could not be retrieved through any open-access channel (reported explicitly in the PRISMA flow), so paywalled surveys—particularly from the early 2000s—are underrepresented. Notably, no survey published before 2009 satisfied all criteria, so the 1999–2008 portion of the window is represented in the narrative (Section 3.2) rather than the catalog. The synthesis is therefore intended as a structured map of the field rather than an exhaustive account of every primary system.

Conclusion validity. The challenge taxonomy and the consensus/divergence synthesis (Table 8) involve interpretive judgment. These are grounded in explicit per-challenge evidence and cross-referenced to the surveyed domains in Table 7 to keep the reasoning traceable.

5. Discussion—Challenges and Research Directions

This section presents the core contribution of this meta-survey: a comprehensive, structured analysis of the challenges and research directions identified across the 85 surveyed papers. Following the dual-theme approach of Saeed and Omlin [58], the discussion is organized into two complementary themes:

1. **Theme 1: Cross-cutting challenges** (Section 5.1)—eighteen recurring challenges that transcend individual system categories.




Figure 9: High-level taxonomy of the cloud database meta-survey, organized into two complementary themes.




Figure 10: The eighteen cross-cutting challenges, organized into four clusters.

2. **Theme 2: Lifecycle-phase challenges** (Section 5.2)—challenges organized by the five phases of the cloud database lifecycle.

Figure 9 provides a high-level overview of the dual-theme taxonomy.

5.1. Theme 1: Cross-cutting challenges

Eighteen cross-cutting challenges that recur across multiple cloud database survey domains are identified and organized, in Figure 10, into four clusters—*foundational*, *data-architecture*, *operational*, and *emerging* concerns. Each challenge is discussed as a short narrative that frames the problem, summarizes the current state reported across the surveyed papers (including points of consensus and divergence), and closes with a brief statement of the open problems and research directions—highlighting data- and knowledge-engineering (DKE) relevance—that the surveys leave unresolved.

Three consolidated views support the discussion. Table 7 maps each challenge to the survey domains that address it, revealing coverage density and thematic clusters; Table 8 synthesizes the key points of agreement and divergence across surveyed papers; and Table 9 shows how attention to each challenge has evolved over the study window. Detailed discussion follows in Sections 5.1.1–5.1.18.

Table 7: Cross-cutting challenge coverage across consolidated survey domains. Symbols: • = primary focus, ◦ = substantive discussion, · = mentioned. The rightmost column (n) counts domains addressing each challenge.

#	Challenge	Cloud-Native	HTAP	NoSQL / NewSQL	AI4DB / Learned	Vector / Graph / Multi-Model	Lakehouse	Time-Series	Edge / Fog	NL / LLM-DB	Security / Privacy	Benchmarking	Migration	Distributed DBs	DBaaS	n
<i>Foundational</i>																
C1	Scalability & elasticity	•	◦	•	·		◦	◦	◦			◦		◦	◦	10
C2	Consistency & CAP/PACELC	•	•	•		·			◦					•		6
C3	Performance & benchmarking	•	◦	◦	•		◦					•		◦		7
C4	Security & privacy	◦		·					◦		•		·		◦	6
C5	Cost modeling & optimization	◦					◦					•	◦		◦	5
<i>Data Architecture</i>																
C6	Multi-tenancy & isolation	•									◦	◦			•	4
C7	Data model diversity		·	•		•	◦	◦								5
C8	Heterogeneous query processing	◦	•	·	◦	◦	•			◦				◦		8
C9	Storage disaggregation	•	◦				•									4
C10	Distributed transactions	•	•	•		·			·					•		5
<i>Operations</i>																
C11	Fault tolerance & HA	•		◦					◦					•		4
C12	Migration & portability	◦		·									•		◦	4
C13	Edge-cloud continuum	·							•		·					3
C14	AI-database convergence	◦			•	·	◦			•						5
C15	Serverless & autonomy	•			◦							·			◦	4
<i>Emerging</i>																
C16	Governance & lineage	·					◦			·	•		·		·	5
C17	Sustainability	·			·							◦				3

Continued on next page

Table 7 continued

#	Challenge	Cloud-Native HTAP	NoSQL/ NewSQL	AI4DB / Learned	Vector / Graph / Multi-Model	Lakehouse	Time-Series	Edge / Fog	NL / LLM-DB	Security / Privacy	Benchmarking	Migration	Distributed DBs	DBaaS	<i>n</i>
C18	Standardization & interoperability	◦	◦		•	◦			•			◦	•		7

Table 8: Synthesis of cross-cutting challenges: key similarities (consensus) and differences (divergences) across the 85 surveyed papers, referenced at survey-domain granularity (domain sizes as in Figure 7).

#	Challenge	Survey domains (<i>n</i>)	Key Similarities (Consensus)	Key Differences (Divergences)
C1	Scalability & elasticity	44 surveys in Cloud-Native, HTAP, NoSQL/NewSQL, Lakehouse, Time-Series, Edge/Fog, Benchmarking, Distributed DBs, DBaaS	All surveys agree that disaggregated storage enables elastic scaling; horizontal scale-out is preferred over vertical scaling; auto-scaling is a baseline cloud capability.	Cloud-native surveys focus on component-level elasticity; NoSQL surveys emphasize partition-based data scaling; HTAP surveys and edge/fog surveys stress elasticity under mixed and intermittently connected workloads, respectively.
C2	Consistency & CAP/PACELC	30 surveys in Cloud-Native, HTAP, NoSQL/NewSQL, Edge/Fog, Distributed DBs	CAP/PACELC trade-offs are universally acknowledged; all surveys recognize a spectrum from strong to eventual consistency; configurable consistency is a key cloud differentiator.	NoSQL surveys accept eventual consistency as pragmatic; HTAP surveys insist on strong consistency for transactional correctness; cloud-native surveys present tunable consistency as an application-level choice.
C3	Performance & benchmarking	45 surveys in Cloud-Native, HTAP, NoSQL/NewSQL, AI4DB/Learned, Lakehouse, Benchmarking, Distributed DBs	Cloud-specific factors (network latency, multi-tenancy, cold starts) complicate optimization; existing benchmarks (TPC-C/H, YCSB) are widely considered inadequate for cloud settings.	AI4DB surveys emphasize learned optimizers and indexes; benchmarking surveys focus on measurement methodology and fairness; cloud-native surveys focus on disaggregation overhead.
C4	Security & privacy	17 surveys in Cloud-Native, Edge/Fog, Security/Privacy, DBaaS	Encryption at rest and in transit is a baseline requirement; the shared responsibility model is universally acknowledged; GDPR compliance is a common concern.	Security surveys deeply analyze threat models and advanced techniques (homomorphic encryption, confidential computing); cloud-native surveys treat security as one concern among many; edge/fog surveys emphasize privacy-preserving aggregation.
C5	Cost modeling & optimization	15 surveys in Cloud-Native, Lakehouse, Benchmarking, Migration, DBaaS	Cost is a primary driver of cloud adoption; pricing complexity (10–100× variation) is widely noted; pay-per-use models are a defining benefit.	Benchmarking surveys focus on total cost of ownership comparison; lakehouse surveys weigh storage–compute cost trade-offs; migration surveys highlight hidden migration and lock-in costs; cloud-native surveys note multi-dimensional pricing complexity.

Continued on next page

Table 8 continued

#	Challenge	Survey domains (n)	Key Similarities (Consensus)	Key Differences (Divergences)
C6	Multi-tenancy & isolation	19 surveys in Cloud-Native, Security/Privacy, Benchmarking, DBaaS	Noisy-neighbor performance interference is universally recognized; the trade-off between resource sharing efficiency and isolation is a fundamental tension.	Cloud-native surveys focus on architectural isolation levels (shared-nothing to shared-table); security surveys emphasize tenant data isolation; benchmarking surveys measure interference under controlled conditions.
C7	Data model diversity	41 surveys in NoSQL/NewSQL, Vector/Graph/Multi-Model, Lakehouse, Time-Series	Polyglot persistence is the operational reality; cross-model query support is needed; multi-model databases are gaining traction.	NoSQL surveys favor model-specific optimization; multi-model surveys advocate unified backends; specialized surveys argue that model-native operations (vector, graph, time-series) cannot be efficiently reduced to a single backend.
C8	Heterogeneous query processing	44 surveys in Cloud-Native, HTAP, AI4DB/Learned, Vector/Graph/Multi-Model, Lakehouse, NL/LLM-DB, Distributed DBs	Federated query processing across engines is increasingly necessary; predicate pushdown is the baseline optimization technique.	HTAP surveys focus on OLTP-OLAP workload routing; lakehouse surveys emphasize open-format access layers; multi-model surveys and natural-language interface surveys address cross-model and intent-driven query routing.
C9	Storage disaggregation	26 surveys in Cloud-Native, HTAP, Lakehouse, NoSQL/NewSQL	Compute-storage separation is the defining cloud-native architectural principle; write amplification reduction is a key benefit; independent layer scaling is a primary motivation.	Cloud-native surveys focus on log-as-database architectures (Aurora, PolarDB); lakehouse surveys emphasize open table formats; HTAP surveys address separate analytical and transactional storage tiers.
C10	Distributed transactions	28 surveys in Cloud-Native, HTAP, NoSQL/NewSQL, Distributed DBs	2PC and consensus protocols (Paxos/Raft) are the dominant approaches; the consistency-latency trade-off is universally acknowledged.	Cloud-native surveys cite clock synchronization innovations (TrueTime, HLC); NoSQL surveys accept weaker guarantees (CRDTs, eventual consistency); HTAP surveys focus on OLTP-OLAP transaction isolation.
C11	Fault tolerance & HA	27 surveys in Cloud-Native, NoSQL/NewSQL, Edge/Fog, Distributed DBs	Multi-AZ replication is the baseline approach; automated failover is expected; RTO/RPO targets drive architectural decisions.	Cloud-native surveys focus on quorum protocols (e.g., Aurora's 4/6 write quorum); NoSQL surveys emphasize tunable replication factors; edge/fog surveys must handle intermittent connectivity.

Continued on next page

Table 8 continued

#	Challenge	Survey domains (<i>n</i>)	Key Similarities (Consensus)	Key Differences (Divergences)
C12	Migration & portability	8 surveys in Cloud-Native, Migration, DBaaS	Migration is complex, error-prone, and labor-intensive; schema heterogeneity is the primary technical barrier; vendor lock-in is widely criticized.	Migration surveys focus on methodology and zero-downtime techniques; cloud-native surveys treat migration as a cloud adoption path; NoSQL surveys highlight data model mismatch challenges.
C13	Edge–cloud continuum	2 surveys in Edge/Fog	Data placement across edge–fog–cloud tiers requires intelligent policies; intermittent connectivity is the defining constraint.	Edge/fog surveys focus on resource-constrained processing and synchronization; cloud-native surveys view edge as a cloud extension; security surveys emphasize privacy-preserving edge aggregation.
C14	AI–database convergence	18 surveys in Cloud-Native, AI4DB/Learned, Lakehouse, NL/LLM–DB	Both AI4DB and DB4AI directions are growing rapidly; production deployment of learned components remains limited; training data management is a shared challenge.	AI4DB surveys focus on individual learned components (indexes, optimizers); natural-language and retrieval-augmented surveys address LLM-driven querying and retrieval; cloud-native surveys view AI as a system enhancement.
C15	Serverless & autonomy	17 surveys in Cloud-Native, AI4DB/Learned, DBaaS	Serverless abstracts infrastructure management; cold-start latency is a known limitation; pay-per-query suits variable workloads.	Cloud-native surveys treat serverless as one deployment model and catalog autoscaling mechanisms; AI4DB surveys envision full-stack autonomous databases that self-tune beyond the serverless abstraction.
C16	Governance & lineage	11 surveys in Lakehouse, Security/Privacy	End-to-end lineage tracking is increasingly important; regulatory compliance (GDPR, CCPA) drives governance adoption; metadata management is fragmented across tools.	Security surveys focus on access control and audit logging; lakehouse surveys emphasize unified data cataloging; migration surveys track lineage across system transitions; cloud-native and NL-interface surveys note metadata’s role in querying.
C17	Sustainability	4 surveys in Benchmarking	Energy-proportional computing is an aspirational goal; carbon-aware data placement is an emerging research direction.	Cloud-native surveys note the energy implications of disaggregation; AI4DB surveys flag the training cost of learned components; benchmarking surveys propose adding energy and carbon metrics; most other survey domains do not yet address sustainability.
C18	Standardization & interoperability	42 surveys in Cloud-Native, NoSQL/NewSQL, Vector/Graph/Multi-Model, Lakehouse, Migration	Vendor lock-in is widely criticized; SQL remains the primary query lingua franca; open formats (Iceberg, Delta Lake, Hudi) are gaining traction.	Graph and specialized surveys highlight query-language fragmentation (Cypher, SPARQL, GQL); NoSQL surveys focus on API heterogeneity; multi-model and natural-language interface surveys push unified access, while migration surveys stress portable formats.

Table 9: Temporal attention to each cross-cutting challenge: the number of surveyed papers, by publication period, whose primary domain substantively (●/○) addresses the challenge. Cell shading is scaled within each row; attention concentrates in 2023–2026, mirroring the overall publication trend (Figure 6).

#	Challenge	2009–2015	2016–2019	2020–2022	2023–2024	2025–2026	Total
C1	Scalability & elasticity	3	10	11	15	5	44
C2	Consistency & CAP/PACELC	2	6	9	10	3	30
C3	Performance & benchmarking	2	7	12	16	8	45
C4	Security & privacy	3	5	3	4	2	17
C5	Cost modeling & optimization	1	4	2	6	2	15
C6	Multi-tenancy & isolation	3	5	3	5	3	19
C7	Data model diversity	3	9	10	14	5	41
C8	Heterogeneous query processing	1	11	6	19	7	44
C9	Storage disaggregation	2	3	9	9	3	26
C10	Distributed transactions	2	6	8	9	3	28
C11	Fault tolerance & HA	2	6	8	8	3	27
C12	Migration & portability	1	4		2	1	8
C13	Edge–cloud continuum			1	1		2
C14	AI–database convergence		3	3	7	5	18
C15	Serverless & autonomy	1	4	2	5	5	17
C16	Governance & lineage	2	2	3	3	1	11
C17	Sustainability			1	2	1	4
C18	Standardization & interop.	3	9	10	15	5	42
Total (challenge-mentions)		31	94	101	150	62	438

5.1.1. Scalability and elasticity

Scalability—handling growing workloads by adding resources—and elasticity—provisioning and releasing resources in response to demand—are the foundational promises of the cloud, and every surveyed cloud-native and NoSQL review treats them as first-order concerns [6, 11]; efforts to quantify elasticity empirically date back to the first wave of cloud data stores [34]. In practice, cloud-native engines such as Aurora [2] and Snowflake [3] achieve storage elasticity through compute-storage disaggregation, while compute elasticity is delivered by auto-scaling and replication [60]; serverless offerings push this further by scaling to zero between requests [27]. The surveyed literature nonetheless converges on a shared limitation: today’s elasticity is coarse-grained, operating at the level of whole instances rather than individual components, so that *fine-grained* elasticity—independently scaling a single query operator, the transaction manager, or the log service—remains largely unrealized [6].

Several obstacles recur across the reviews. Cold-start latency undermines the responsiveness that makes scale-to-zero attractive [27]; state migration during scale-out and scale-in introduces transient performance degradation; and predictive autoscaling depends on workload-forecasting models that are still maturing [61]. Cross-component elasticity in fully disaggregated architectures is the least understood of all, because the interactions among independently scaled layers are difficult to model.

Open problems and research directions. The central open problem is a principled account of *elastic data architectures* that links workload patterns, resource allocation, performance guarantees, and cost. From a data- and knowledge-engineering standpoint, ontological models of cloud-database configurations and their scaling behavior could turn today’s heuristic autoscaling into reasoning over an explicit model, enabling predictable, component-level elasticity.

5.1.2. Consistency, availability, and partition tolerance revisited

The CAP theorem [62] and its PACELC refinement [63] have shaped distributed database design for over a decade, but the surveyed work shows that geo-distribution, multi-region deployment, and the proliferation of configurable consistency levels give these trade-offs new dimensions in the cloud [64]. Cloud databases now expose a spectrum of guarantees: Spanner provides external consistency via TrueTime [46], whereas systems such as Cosmos DB offer several intermediate levels from strong to eventual. Surveys of NoSQL stores [11, 12] and of HTAP systems [7] consistently single out consistency management as a defining difficulty, and they disagree on how much weakening is acceptable—NoSQL reviews accept eventual consistency as pragmatic, while HTAP and NewSQL reviews insist on strong guarantees for transactional correctness.

What the surveys leave open is largely a matter of reasoning and tooling. Developers struggle to predict how a given consistency level affects application correctness [64]; adaptive consistency that adjusts to workload requirements is frequently proposed but rarely deployed; and cross-model consistency—reconciling, say, document and graph views of the same data—is barely studied in multi-model settings [52].

Open problems and research directions. Bridging the gap between formal consistency models and everyday development calls for knowledge-based specifications of consistency *contracts*—machine-checkable statements of the guarantees a service offers and an application requires—so that compatibility can be verified rather than assumed, a natural fit for formal-methods and DKE techniques.

5.1.3. Performance optimization and benchmarking

Performance optimization in the cloud spans query optimization, indexing, caching, resource allocation, and workload management, and benchmarking supplies the empirical ground truth against which systems are compared [65]. The AI4DB and cloud-native reviews document substantial progress on the optimization side—learned query optimization [66], learned indexes [67], and automated knob tuning [68]—yet they emphasize that cloud-specific factors such as network latency, disaggregation overhead, multi-tenant interference, and cold starts confound techniques designed for single-node systems [9, 6]. On the measurement side, the benchmarking reviews are unanimous that classical benchmarks (TPC-C, TPC-H, YCSB) fail to capture cloud realities such as elasticity, pay-per-use pricing, and variable network performance [65, 69, 70, 71].

The unresolved issues follow directly. Cloud-aware benchmarks that incorporate elasticity, multi-tenancy, and pricing are still lacking; fair cross-provider comparison is confounded by differing hardware, pricing, and service-level agreements; and end-to-end performance in disaggregated stacks emerges from complex interactions among independently scaled layers that no current benchmark isolates.

Open problems and research directions. A promising direction is a *benchmark knowledge base* that formally records configurations, workload characteristics, hardware, and results in a structured, queryable representation; making benchmark provenance explicit would render cloud-database evaluations reproducible and genuinely comparable—an evaluation problem with a strong knowledge-engineering core.

5.1.4. Security and privacy

Relative to on-premises systems, cloud databases face an enlarged attack surface arising from shared infrastructure, network-based access, and the need to trust the provider [72, 73, 74]. The security-focused surveys catalog the established defenses—encryption at rest and in transit, access control, audit logging, and vulnerability management [75, 30, 76, 77]—and the more advanced techniques that target computation over protected data, including homomorphic encryption for querying ciphertext [78, 79], architectures giving clients distributed, concurrent access to encrypted cloud databases without a trusted broker [80], private-database models for outsourced data [81], secure multi-party computation, and confidential computing in hardware enclaves. The reviews differ in emphasis: dedicated security surveys analyze threat models in depth, whereas cloud-native reviews treat security as one concern among many and edge-oriented work stresses privacy-preserving aggregation close to the data source.

Across the literature the same hard problems persist: the performance overhead of encryption remains prohibitive for analytics; key management across multi-cloud and hybrid deployments is error-prone; the “honest but curious” provider is difficult to defend against without heavy penalties; and reconciling GDPR’s right to erasure with analytic utility is unsolved.

Open problems and research directions. *Privacy ontologies* that formally model data sensitivity, access policies, and regulatory obligations would let compliance be checked and enforced automatically across heterogeneous cloud deployments, turning today’s manual, document-driven compliance into machine-reasoned policy—squarely a DKE contribution.

5.1.5. Cost modeling and optimization

Cloud databases bill along many dimensions at once—compute hours, storage, I/O operations, network egress, backups, and premium features such as encryption or multi-region replication—which makes cost both a primary adoption driver and a persistent source of surprise [25]. Empirical studies show that the cost of the same workload can vary by one to two orders of magnitude depending on configuration [25], and although work on serverless analytics [27] and the cloud-native reviews [6] flag cost as critical, they offer little quantitative guidance. The benchmarking and migration reviews add the complementary observations that total-cost-of-ownership comparison and hidden lock-in costs are rarely modeled rigorously [69, 82].

Consequently, pre-execution cost estimation for complex analytical queries remains unreliable; multi-objective optimization over cost, performance, and consistency lacks mature frameworks; the trade-offs among reserved, on-demand, and spot capacity for database workloads are understudied; and vendor lock-in obscures cross-provider comparison.

Open problems and research directions. Formal *cost knowledge models*—ontologies that capture provider pricing structures, workload profiles, and optimization strategies—could support automated, explainable cost optimization, and would give the field a shared vocabulary for reasoning about cloud economics rather than per-provider spreadsheets.

5.1.6. Multi-tenancy and resource isolation

Multi-tenancy lets providers serve many customers from shared infrastructure, improving utilization and lowering cost, at the price of harder performance isolation, data separation, and fair resource allocation [6]. The surveyed systems realize tenancy at several granularities—shared-nothing (an instance per tenant), shared-process (logical isolation within one engine), and shared-table (row-level security over common tables)—and the reviews repeatedly identify “noisy neighbor” interference, in which one tenant’s load degrades another’s latency, as the dominant operational hazard [83].

Open problems and research directions. Achieving strong performance isolation without dedicating hardware remains difficult; workload-aware scheduling that adapts to heterogeneous tenants is immature; and the fundamental tension between sharing efficiency and isolation guarantees still lacks a principled, cost-aware resolution. Quantitative models that predict interference from tenant workload signatures are a promising avenue.

5.1.7. *Data model diversity and multi-model integration*

The cloud has multiplied the data models in routine use—relational, document, graph, key-value, wide-column, vector, and time-series—and applications increasingly need to query across several at once [52, 84, 85, 86]. Multi-model engines such as Cosmos DB, ArangoDB, and SurrealDB expose multiple model interfaces over a single backend, and the NoSQL and multi-model surveys report growing adoption alongside a candid account of the limits: cross-model query optimization and unified transaction semantics remain weak [11, 87, 88, 89, 52]. The reviews also diverge on strategy—NoSQL surveys favor model-specific optimization, multi-model surveys advocate unified backends, and specialized reviews argue that vector, graph, and time-series operations cannot be reduced efficiently to a single engine.

Open problems and research directions. The persistent gaps are unified query languages spanning relational, graph, document, and vector operations; cross-model optimization of queries that combine, say, graph traversal with relational joins and similarity search; and schema mapping across models. Using ontologies to provide a single conceptual layer over diverse physical models—*knowledge representation for multi-model data*—is a direct DKE opportunity.

5.1.8. *Query processing across heterogeneous engines*

Modern cloud data stacks rarely rely on a single engine; OLTP, OLAP, graph, search, and vector systems must be queried in concert, making federated query processing and data virtualization essential [7, 90]. Engines such as Presto/Trino, BigQuery Omni, and AWS Lake Formation already provide federation, and the HTAP reviews analyze how queries should be routed to the appropriate engine [7]; yet they agree that cross-engine optimization is still rudimentary, seldom going beyond predicate pushdown. The unresolved problems are cost-based optimization across engines with incompatible cost models and capabilities, preserving transactional guarantees across engine boundaries, and the semantic understanding of data needed to route queries intelligently rather than syntactically.

Open problems and research directions. *Ontology-mediated query answering* [23] offers a principled basis for posing a single query against many heterogeneous engines through a shared conceptual schema, connecting cloud query federation to a long line of DKE research on data integration.

5.1.9. *Storage disaggregation and log-as-database*

Separating compute from storage into independently scalable layers is the defining architectural move of the cloud-native era, and the “log-as-database” pattern pioneered by Aurora—treating the redo log itself as the system of record—pushes intelligence into the storage tier [2, 6]. This design reduced write amplification and enabled fast failover, and systems such as PolarDB [17] and Socrates [4] carry disaggregation further by splitting out additional components. The cloud-native surveys report that while disaggregation simplifies operations and scaling, it introduces its own difficulties in remote-I/O latency, cache coherence, and log management [6].

Open problems and research directions. Outstanding issues include masking the latency gap between remote storage and local SSDs, coordinating caches across compute and storage layers, and managing the log at scale through compaction, garbage collection, and efficient replay. Cost- and latency-aware models of where computation should run within a disaggregated stack are a promising research direction.

5.1.10. *Transaction processing in distributed cloud environments*

Distributed transaction processing must balance ACID guarantees against performance and scalability, and the surveyed systems span the full design space from strongly consistent protocols (two-phase commit, Paxos/Raft) to relaxed models such as CRDTs and eventual consistency [46, 16]. Spanner demonstrated that globally distributed, strongly consistent transactions are feasible using synchronized clocks [46], and CockroachDB and YugabyteDB provide comparable guarantees with hybrid logical clocks [16]; the NewSQL and HTAP reviews nonetheless report that sustaining strong consistency at high throughput for mixed workloads is still hard [7].

Open problems and research directions. The recurring open problems are reducing commit latency for geo-distributed transactions without weakening consistency, resolving conflicts efficiently across regions, and isolating transactional and analytical work within a single HTAP engine so that neither degrades the other.

5.1.11. *Fault tolerance, recovery, and high availability*

Cloud databases must survive hardware faults, network partitions, availability-zone outages, and even whole-region failures while preserving durability and minimizing downtime [2, 6]. The prevailing approach is multi-availability-zone replication—typically three-way—with automated failover, exemplified by Aurora’s six-replica, three-zone write

quorum [2]. The surveys report that recovery-time objectives of seconds and near-zero recovery-point objectives are attainable within a single region, but remain difficult to guarantee across regions.

Open problems and research directions. Multi-region disaster recovery with near-zero RPO and RTO is still open, as is consistent recovery in disaggregated architectures whose state is scattered across components; the surveyed work also notes the absence of chaos-engineering and fault-injection frameworks tailored to cloud databases, which would let resilience be tested systematically rather than anecdotally.

5.1.12. *Data migration and cross-cloud portability*

Organizations routinely move databases between on-premises and cloud, between providers, and between services, and the migration survey documents how hard it is to preserve consistency, minimize downtime, and transform schemas along the way [82]. Although every major provider offers a migration service, the reviewed evidence is that migrations remain error-prone and labor-intensive, with schema and data-model mismatches the principal technical barrier and vendor lock-in—through proprietary APIs, formats, and features—compounding the difficulty of any subsequent move [91].

Open problems and research directions. The field still lacks zero-downtime migration for large, continuously written databases, automated schema mapping between heterogeneous systems, and portability layers that abstract provider-specific features. *Knowledge-based schema mapping*, which represents schema semantics ontologically to automate translation, ties cloud migration directly to DKE's long-standing strengths in schema integration and data exchange.

5.1.13. *Edge–cloud data continuum*

The spread of IoT devices, autonomous vehicles, and smart-city infrastructure has created demand for database capability at the network edge operating in concert with the cloud [13], a pressure recognized early in evaluations of cloud databases for IoT data [92]. The edge-and-fog survey maps architectures that distribute storage and query processing across edge devices, fog nodes, and cloud backends, and notes that systems such as Azure SQL Edge, AWS Outposts, and locality-aware features in CockroachDB address only parts of the problem [13].

Open problems and research directions. Maintaining consistency across the edge–fog–cloud continuum under intermittent connectivity, deciding intelligently which data to cache, replicate, or move, processing queries under tight resource constraints on edge devices, and aggregating data privately before it reaches the cloud all remain open. Policy models that reason about placement against latency, privacy, and connectivity constraints are a promising direction.

5.1.14. *AI–database convergence (AI4DB and DB4AI)*

The convergence of AI and databases runs in two directions—AI4DB, which uses learning to optimize the database, and DB4AI, which uses database technology to serve AI workloads [9]. The AI4DB surveys document concrete progress in learned query optimization [66], learned indexes [67, 93, 94], automated knob and buffer tuning [68, 95], learned partitioning advisors [96], cardinality estimation, and workload forecasting [61, 97], while DB4AI spans in-database machine learning and optimized data serving for training pipelines. The reviews agree that learned components are rarely yet trusted in production, and that managing their training data is a shared, unsolved concern. The persistent obstacles are deploying learned components with safety guarantees so that they degrade gracefully when a model fails, collecting and curating training data from live workloads, transferring learned behavior across instances and hardware, and—perhaps most important for adoption—explaining why a learned optimizer made a given decision.

Open problems and research directions. *Knowledge engineering for self-driving databases*, which couples explicit domain knowledge about database internals with learned models, is a promising route to autonomous systems that are both robust and explainable—an agenda where DKE and the database community meet directly.

5.1.15. *Serverless and autonomous database operations*

Serverless databases hide infrastructure entirely, provisioning, scaling, patching, and tuning without operator involvement, and autonomous databases go further by using machine learning to self-diagnose and self-heal [45, 27]. Offerings such as Aurora Serverless v2, the Cosmos DB serverless tier, Neon, and PlanetScale realize the serverless interface, while Oracle Autonomous Database is the most ambitious autonomy effort; studies of serverless data systems report that the model suits variable workloads well but struggles with sustained high throughput and cold-start latency [27].

Open problems and research directions. Predictable performance under serverless autoscaling, transparent exposure of the cost–performance trade-off, and full-stack autonomy that unifies self-tuning, self-healing, and self-securing remain open. The cloud-native and AI4DB reviews converge on the view that genuine autonomy will require databases to reason about their own behavior, not merely react—again pointing to a knowledge-grounded approach.

5.1.16. Data governance, provenance, and lineage

As cloud data ecosystems grow, tracking where data came from (provenance), how it was transformed (lineage), and whether its use complies with policy becomes essential for compliance, debugging, and trust [24]. The governance-related surveys identify a common set of needs—metadata management, data cataloging, access control, quality monitoring, and regulatory compliance—and report that cloud-native tools such as AWS Glue Data Catalog, Azure Purview, and Google Dataplex address them only partially, leaving the landscape fragmented and unstandardized. The unresolved problems are end-to-end lineage that spans heterogeneous services and processing engines, automated checking against regulatory requirements that themselves keep changing, and provenance-aware query processing that records how each result was derived.

Open problems and research directions. Representing and linking metadata as a *knowledge graph* over cloud data assets, and reasoning over it, is a natural DKE contribution that would unify cataloging, lineage, and compliance under a single queryable model.

5.1.17. Sustainability and energy efficiency

Data centers consumed an estimated 1–1.5% of global electricity in 2023, with projections of 3–4% by 2030 as AI and cloud workloads grow, and cloud databases—heavy consumers of compute and storage—contribute directly to that footprint [98]. Work on green computing identifies opportunities in energy-efficient query processing, scheduling that follows renewable-energy availability, and carbon-aware data placement; major providers now report carbon metrics, but the surveyed evidence is that database-level energy optimization is still nascent.

Open problems and research directions. Energy-proportional databases that consume resources in line with actual load, carbon-aware placement and replication that favor cleaner grids, and benchmarks that report energy and carbon alongside performance are the principal open directions; sustainability is also the least-addressed challenge across the surveyed domains, marking it as an opportunity rather than a settled topic.

5.1.18. Standardization and interoperability

The proliferation of cloud database systems, query languages, APIs, and formats has produced an interoperability problem: SQL remains the lingua franca for relational data, but graph databases alone span Cypher, SPARQL, Gremlin, and GQL, and NoSQL systems expose largely proprietary APIs [8]. The surveys note genuine progress—GQL’s emergence as an ISO standard, the SQL/PGQ extension for property graphs, and the consolidation of open table formats such as Iceberg, Delta Lake, and Hudi—while cautioning that vendors continue to differentiate through proprietary extensions, keeping true interoperability out of reach.

Open problems and research directions. Universal query interfaces spanning relational, graph, document, and vector operations, standardized control-plane APIs for provisioning and monitoring, and open formats and protocols that blunt lock-in are the recurring asks. Because they are fundamentally about shared schemas and vocabularies, these are interoperability problems to which knowledge-engineering methods apply directly.

5.2. Theme 2: Challenges by cloud database lifecycle phases

Complementing the cross-cutting perspective, challenges are organized by the five phases of the cloud database lifecycle: design, deployment, operation, optimization, and evolution. Figure 11 illustrates this lifecycle with key challenges mapped to each phase, and Table 10 groups the surveyed papers by the phase(s) they address. The distribution is itself revealing: design and optimization are the most heavily surveyed phases, whereas the *evolution* phase—schema change, technology migration, and the adoption of new paradigms such as vector and lakehouse systems—is addressed by the fewest surveys, marking it as the least mature stage of the lifecycle and a clear opportunity for future work.

5.2.1. Design phase

The design phase encompasses decisions made before a cloud database is deployed, including schema design, data model selection, workload characterization, and compliance architecture.




Figure 11: The cloud database lifecycle with key challenges mapped to each phase.

Table 10

Surveyed papers grouped by the lifecycle phase(s) they address (derived from the catalog in Table 11; a survey may span several phases, so the counts sum to more than 85). Representative examples are cited by S-ID; the complete per-phase assignment is part of the archived catalog data.

Phase	Representative challenges	Examples	<i>n</i>
Design	Schema design, data-model selection, workload characterization, compliance-by-design	S5, S11, S47, S64, S82	54
Deployment	Provider selection & lock-in, configuration automation, data ingestion/migration, multi-region setup	S4, S9, S18, S57, S84	22
Operation	Monitoring & observability, backup/recovery, security operations, HTAP workload management	S1, S5, S13, S40, S72	39
Optimization	Learned query optimization, index/materialized-view selection, knob tuning, cost-performance balance	S2, S21, S32, S54, S78	65
Evolution	Schema evolution & online DDL, technology migration, new-paradigm adoption, LLM integration	S4, S9, S34, S44, S60	6

Schema design for cloud-native environments. Traditional schema design principles (normalization, ER modeling) must be adapted for cloud databases where denormalization may improve performance due to the cost of distributed joins [6]. Document and wide-column stores require fundamentally different schema design approaches centered on access patterns rather than entity relationships [11, 99, 100, 101].

Workload characterization. Accurate workload modeling is essential for selecting the right cloud database service, configuring resources, and estimating costs. Surveys on AI4DB [9] discuss workload forecasting using ML, but applying these techniques during the design phase—before production data exists—remains challenging.

Data model selection. Choosing between relational, document, graph, key-value, wide-column, vector, or multi-model databases requires understanding workload characteristics, query patterns, consistency requirements, and cost constraints. No systematic decision framework exists that accounts for all these dimensions [52].

Compliance-by-design. Incorporating regulatory requirements (GDPR, HIPAA, data sovereignty) into the initial database design rather than retrofitting compliance is increasingly recognized as essential [29]. This includes data residency constraints, encryption requirements, and audit trail design.

5.2.2. Deployment phase

The deployment phase covers provider selection, infrastructure configuration, data ingestion, and multi-region setup.

Provider selection and vendor lock-in. Choosing a cloud database provider involves evaluating technical capabilities, pricing, compliance certifications, geographic coverage, and ecosystem integration. Vendor lock-in—through proprietary features, APIs, and data formats—is consistently identified as a top concern across surveys [91, 6].

Configuration automation. Infrastructure-as-Code (IaC) tools (Terraform, Pulumi, CloudFormation) enable automated database provisioning, but optimal configuration selection—instance types, storage tiers, replication settings, parameter tuning—often requires deep expertise that is difficult to automate [68].

Data ingestion and migration. Initial data loading and ongoing ingestion from diverse sources (operational databases, streaming platforms, file systems) into cloud databases involves ETL/ELT pipeline design, schema mapping, and data quality assurance [82].

Multi-region deployment. Deploying databases across multiple geographic regions for latency, availability, and compliance reasons introduces challenges in data synchronization, conflict resolution, and region-aware routing [46].

5.2.3. Operation phase

The operation phase covers the day-to-day management of running cloud database systems.

Monitoring and observability. Cloud databases generate vast amounts of telemetry data (metrics, logs, traces). Effective monitoring requires understanding normal behavior patterns, detecting anomalies, and correlating events across distributed components [61]. Work on self-driving databases [45] emphasizes the need for ML-driven anomaly detection and root cause analysis, and production experience with diagnosing intermittent slow queries at cloud scale shows how difficult root-cause attribution remains [102].

Backup, recovery, and disaster recovery. While cloud providers automate basic backups (point-in-time recovery, snapshots), disaster recovery across regions, testing recovery procedures, and meeting stringent RPO/RTO targets remain operational challenges [2].

Security operations. Ongoing security management includes access control auditing, encryption key rotation, vulnerability patching, network security configuration, and compliance monitoring. The dynamic nature of cloud environments (auto-scaling, serverless) complicates security posture management [75].

Real-time workload management (HTAP). For HTAP systems [7], operational challenges include managing resource contention between transactional and analytical workloads, maintaining freshness guarantees for analytical queries, and preventing analytical queries from degrading OLTP performance.

5.2.4. Optimization phase

The optimization phase encompasses techniques for improving performance and efficiency of running cloud databases.

Learned query optimization. ML-based query optimizers [66, 9] that learn from workload history to generate better query plans represent a paradigm shift from traditional cost-based optimization. However, deploying learned optimizers in production requires addressing safety (avoiding catastrophic regressions), training efficiency, and adaptability to workload changes.

Automated index and materialized view selection. Cloud databases increasingly offer automated index recommendation (e.g., Azure SQL automatic tuning, Amazon Redshift automatic materialized views), but surveys [9] report that these systems are often conservative, covering only a fraction of possible optimizations.

Knob tuning. Database configuration parameters (buffer pool size, parallelism degree, checkpoint intervals, etc.) significantly impact performance. ML-based tuning systems [68, 95] can explore the high-dimensional configuration space more effectively than manual tuning, but transferring tuning knowledge across different workloads and hardware remains an open problem.

Cost–performance trade-off optimization. In cloud environments, performance optimization cannot be divorced from cost. The optimal configuration minimizes cost subject to performance constraints (or maximizes performance subject to budget constraints), requiring multi-objective optimization [25].

5.2.5. Evolution phase

The evolution phase addresses long-term changes in database systems, technologies, and requirements.

Schema evolution and online DDL. Changing database schemas in production without downtime (online DDL) is a critical capability for agile development. Cloud-native databases have improved online DDL support, but complex migrations (column type changes, table splits, data backfills) remain challenging [6].

Technology migration. Migrating from one database technology to another (e.g., from a relational database to a document store, or from a legacy system to a cloud-native database) is a major undertaking that surveys identify as poorly supported by existing tools [82].

Adapting to new paradigms. The emergence of vector databases for AI workloads [5], the lakehouse paradigm [15], and serverless computing [27] requires existing systems to evolve or be replaced. Understanding when and how to adopt new paradigms is a strategic challenge.

LLM integration. The rapid rise of large language models (LLMs) is creating new requirements for cloud databases, from vector storage and retrieval for RAG pipelines to natural language interfaces for database querying and administration (discussed in detail in Section 6).

6. Cloud Databases and Large Language Models—A Convergence Case Study

The convergence of cloud databases and large language models (LLMs) represents one of the most dynamic and rapidly evolving areas in modern data management. This section examines this convergence as an illustrative case study, paralleling Saeed and Omlin's [58] treatment of XAI in medicine as a domain-specific application.

6.1. LLM-augmented database interfaces

Natural language interfaces to databases (NLIDB) have been a long-standing research goal [103]. The advent of LLMs has dramatically improved text-to-SQL accuracy, with systems like DIN-SQL [104], DAIL-SQL [105], and MAC-SQL [106] achieving over 85% execution accuracy on the Spider benchmark [107].

Current capabilities. LLM-based text-to-SQL systems leverage schema-aware prompting, few-shot in-context learning, and chain-of-thought reasoning to translate natural language queries into SQL. Cloud vendors are integrating these capabilities: Amazon Q for databases, Azure AI integration with SQL, and Google Duet AI for BigQuery.

Open challenges.

- Complex queries involving multiple joins, subqueries, and window functions remain error-prone.
- Schema understanding for large databases with hundreds of tables and complex relationships.
- Multi-turn conversational interfaces that maintain context across query refinements.
- Handling ambiguity, domain-specific terminology, and user intent clarification.
- Security risks from SQL injection through natural language prompts.

6.2. Vector databases as LLM infrastructure

Vector databases have emerged as critical infrastructure for LLM applications, particularly retrieval-augmented generation (RAG) pipelines [5].

Architecture and workloads. RAG pipelines embed documents into high-dimensional vectors, store them in vector databases, and retrieve relevant context for LLM prompts based on semantic similarity. This creates unique workload patterns: bulk embedding ingestion, high-throughput approximate nearest neighbor (ANN) search, and hybrid queries combining vector similarity with metadata filtering.

System landscape. The vector database market has rapidly diversified, including purpose-built systems (Pinecone, Weaviate, Milvus, Qdrant, Chroma), vector extensions to existing databases (pgvector for PostgreSQL, Elasticsearch vector search), and cloud-native offerings (Amazon OpenSearch vector engine, Azure AI Search, Google Vertex AI Vector Search) [5].

Open challenges.

- Scalability and freshness—keeping vector indexes up-to-date as source documents change.
- Multi-modal vector search combining text, image, and structured metadata.

- Quality metrics for retrieval that correlate with downstream LLM task performance.
- Cost optimization for storing and querying billions of high-dimensional vectors.
- Integration with existing cloud database ecosystems (transactions, access control, governance).

6.3. LLMs for database administration

LLMs are increasingly applied to automate database administration tasks, including performance diagnosis, query optimization, and anomaly resolution.

Systems and approaches. D-Bot [108] uses LLMs for database diagnosis, performing root cause analysis of performance anomalies by querying system views and analyzing metrics. DB-GPT [109] provides an LLM-powered framework for database interaction, supporting text-to-SQL, data analysis, and administration tasks. GPTuner [110] leverages LLMs for database knob tuning, combining GPT-4 knowledge with Bayesian optimization.

Open challenges.

- Reliability—LLMs can hallucinate incorrect diagnostic conclusions or generate harmful administration commands.
- Safety guardrails—preventing LLM-driven administration from executing destructive operations.
- Integrating LLM reasoning with structured database knowledge (system catalogs, performance models).
- Evaluation frameworks for LLM-based DBA tools that measure real-world effectiveness.

6.4. Challenges and open problems at the convergence

The LLM–database convergence creates several overarching challenges:

1. **Infrastructure co-design.** Jointly optimizing cloud database infrastructure for both traditional workloads and LLM-related workloads (embedding generation, vector search, context retrieval) requires new architectural thinking.
2. **Data governance for AI.** Ensuring that LLM applications respect data access policies, privacy regulations, and audit requirements when retrieving context from cloud databases.
3. **Evaluation and benchmarking.** Standardized benchmarks for vector database performance, text-to-SQL accuracy on real-world schemas, and LLM-based DBA effectiveness are needed.
4. **Cost management.** LLM API calls combined with cloud database queries create complex cost structures that are difficult to predict and optimize.

6.5. What do we think? Open questions for cloud databases and LLMs

Looking across the challenges catalogued in Section 5, the cloud-database–LLM convergence raises a set of questions that the surveyed literature has only begun to address. They are offered not as a checklist but as prompts for future work in this domain—and, by analogy, for other data-intensive domains:

- How should a cloud database co-schedule traditional query workloads and LLM-driven workloads (embedding generation, vector search, context retrieval) on shared, elastic infrastructure without one starving the other?
- When an LLM retrieves context from governed cloud data, how can access policies, privacy regulations, and audit obligations be enforced *through* the retrieval path rather than around it?
- What would a faithful benchmark for this convergence measure—vector-search quality, text-to-SQL accuracy on real enterprise schemas, or the end-to-end utility of an LLM-assisted DBA—and who maintains its ground truth?
- Can the reliability of LLM-based administration be bounded well enough to let it act on production systems, and what guardrails make graceful degradation the default when the model is wrong?

- How should the combined cost of LLM inference and cloud-database execution be modeled and optimized as a single objective, given that today they are billed and tuned independently?
- What role can knowledge graphs and ontologies play in grounding LLM reasoning in the actual schema, constraints, and semantics of a cloud database—turning a plausible answer into a correct one?

These questions show how the cross-cutting and lifecycle challenges of Section 5 resurface, sharpened, when cloud databases meet LLMs; the same analysis could be re-instantiated for other domains such as scientific data management or real-time analytics.

7. Conclusions

This paper has presented the first systematic meta-survey of cloud database research, synthesizing findings from 85 content-verified survey papers identified through a reproducible, scripted search covering January 1999 to July 2026. Applying the Kitchenham–Charters protocol with PRISMA 2020 reporting, the scattered landscape of cloud database challenges has been consolidated into a structured, dual-theme taxonomy.

For the first theme, eighteen cross-cutting challenges are identified and synthesized: (1) scalability and elasticity; (2) consistency and the CAP/PACELC trade-offs; (3) performance optimization and benchmarking; (4) security and privacy; (5) cost modeling and optimization; (6) multi-tenancy and resource isolation; (7) data-model diversity and multi-model integration; (8) heterogeneous query processing; (9) storage disaggregation and the log-as-database pattern; (10) distributed transaction processing; (11) fault tolerance, recovery, and high availability; (12) data migration and cross-cloud portability; (13) the edge–cloud data continuum; (14) AI–database convergence; (15) serverless and autonomous operation; (16) data governance, provenance, and lineage; (17) sustainability and energy efficiency; and (18) standardization and interoperability. Across these, *scalability*, *consistency*, *security*, *cost*, and *AI–database convergence* recur in the majority of the surveyed domains (Table 9), indicating where progress would benefit the field most broadly.

For the second theme, the challenges are mapped to the five phases of the cloud database lifecycle—design, deployment, operation, optimization, and evolution. Design and optimization are the most thoroughly surveyed phases, whereas evolution (schema change, technology migration, and new-paradigm adoption) is the least mature and most under-served (Table 10), making it a particularly promising target for future work.

Two observations cut across both themes. The convergence of cloud databases with large language models is moving rapidly from research to practice through vector databases, text-to-SQL, and LLM-based administration, yet reliability, governance, and evaluation remain unresolved (Section 6). And from a knowledge-engineering standpoint, cloud databases offer rich openings for DKE techniques—knowledge graphs for metadata and lineage, ontology-mediated integration and query answering, and formal models for consistency, cost, and schema evolution.

Accordingly, this meta-survey is intended to serve three purposes: (1) as a comprehensive reference for researchers entering the cloud database space, providing a structured map of the territory and its open problems; (2) as a strategic guide for practitioners navigating technology selection, deployment, and evolution; and (3) as a bridge between the cloud database and the data-and-knowledge-engineering communities.

Limitations. Three limitations qualify these conclusions. First, the corpus is deliberately restricted to surveys that could be content-verified against a resolvable DOI or arXiv identifier, trading breadth for confirmability, so some relevant but harder-to-verify or paywalled surveys may be absent. Second, study selection, quality assessment, and data extraction were performed by a single reviewer (Section 4.5). Third, as a survey of surveys, the findings inherit the currency limitations of the primary surveys, so the fastest-moving topics—most notably LLM–database convergence—may be under-represented relative to their present activity.

Future work

Based on the analysis of 85 surveys and the challenges discussed in Sections 5 and 6, several priority research directions emerge:

- **Foundational systems.** Cloud-native benchmarks that capture elasticity, multi-tenancy, and pricing; the transition of learned optimizers, learned indexes, and ML-based tuning from prototypes to production-ready, safely-degrading components; and a maturing vector-database ecosystem with standardized APIs and retrieval-quality metrics.

- **Architecture and query processing.** Cross-cloud middleware and portable open formats that reduce vendor lock-in; and federated query processing with consistent transactions across heterogeneous engines and the edge–fog–cloud continuum.
- **Intelligence and autonomy.** Autonomous databases whose learned decisions are explainable and grounded in database knowledge; and intent-driven natural-language interfaces that let users express goals directly while preserving correctness and access control.
- **Governance and sustainability.** Privacy-preserving analytics and end-to-end lineage that scale to regulatory needs; and energy- and carbon-aware database operation that converges the relational, graph, document, vector, and streaming models toward unified, sustainable platforms.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Declaration of Generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used a generative AI large-language-model assistant (i) to help draft, edit, and format portions of the manuscript and to prepare its figures and tables, and (ii) to assist in executing the systematic review pipeline—running the scripted literature search, applying the written screening rubrics at the title/abstract and full-text stages, and proposing quality-assessment scores—with every machine-assisted decision logged in the archived screening records (see the Data availability statement). After using this tool, the authors reviewed and edited the content, re-checked all provisional inclusion decisions, adjudicated every borderline case and the venue-integrity audit, and take full responsibility for the content of the publication.

Data availability

This study analyzes previously published survey papers, all of which are cited in the manuscript. The complete identification and selection pipeline is archived with the manuscript sources: the search scripts (`search/search_surveys.py` and companions), the per-query search log (`search_log.json`), the deduplicated candidate pool (`candidates.csv`), the per-record screening-decision log covering every stage of the PRISMA flow (`screening_decisions.csv`), the quality-assessment scores (`qa_scores.csv`), and the frozen machine-readable catalog (`catalog.csv`) from which every figure and table in Section 4.4 is generated. This permits verbatim re-execution of the search and independent audit of every inclusion decision.

A. Complete List of Surveyed Papers with Classification

Table 11 presents the complete catalog of the 85 surveyed papers, classified by primary topic, year, venue, quality-assessment score, and lifecycle phases covered.

Table 11: Master catalog of surveyed papers (S1–S85), ordered chronologically. QA = quality-assessment score (out of 8; inclusion threshold ≥ 6).

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S1	Database as a Service: Challenges and solutions for privacy and... [30]	2009	Security/Privacy	IEEE APSCC	6	Operate
S2	An in-depth comparison of subgraph isomorphism algorithms in gr... [111]	2012	Graph Databases	PVLDB	6.5	Optimize
S3	A Survey of Cloud Database Systems [37]	2014	NoSQL/NewSQL	IEEE IT Professional	6	Design, Operate, Optimize
S4	A Systematic Review of Cloud Computing, Big Data and Databases ... [38]	2014	DBaaS	AMCIS	8	Design, Deploy, Operate, Optimize, Evolve
S5	Confidential database-as-a-service approaches: taxonomy and sur... [31]	2015	Security/Privacy	J. Cloud Comput.	7.5	Design, Operate, Optimize
S6	Choosing the right NoSQL database for the job: a quality attrib... [40]	2015	NoSQL/NewSQL	J. Big Data	7.5	Design, Operate, Optimize
S7	Database-as-a-Service for Big Data: An Overview [112]	2016	DBaaS	IJACSA	7	Design, Deploy, Operate, Optimize
S8	Some problems on graph databases [113]	2016	Graph Databases	Proc. ISP RAS	7	Optimize
S9	Cloud Migration Process: A Survey [82]	2016	Cloud DB Migration	J. Syst. Softw.	7.5	Deploy, Evolve
S10	Foundations of Query Languages for Graph DBs [8]	2017	Graph Databases	ACM Comp. Surv.	8	Design, Optimize
S11	Open Source Time Series Databases: A Survey [50]	2017	Time-Series DBs	BTW	6	Design, Operate

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S12	Enforcing Privacy in Cloud Databases [114]	2017	Security/Privacy	arXiv	6.5	Operate, Optimize
S13	Big data storage technologies: a survey [115]	2017	NoSQL/NewSQL	Front. Inf. Technol. Electron. Eng.	7.5	Design, Operate
S14	Evolutionary Algorithms for Query Op-timization in Distributed ... [116]	2018	AI4DB/Learned Opt.	ADCAIJ	6	Optimize
S15	A Survey on NoSQL Stores [11]	2018	NoSQL/NewSQL	ACM Comp. Surv.	8	Design, Deploy
S16	Saving Large Semantic Data in Cloud: A Survey of the Main DBaaS... [117]	2018	DBaaS	Informatica Economics	6	Design, Deploy
S17	A Survey on Graph Database Management Techniques for Huge Unstr... [118]	2018	Graph Databases	IJECE	7	Design, Optimize
S18	DBaaS Multitenancy, Auto-tuning and SLA Maintenance in Cloud En... [119]	2018	DBaaS	iSys	7	Deploy, Operate, Optimize
S19	Recent NL Interfaces for Databases: A Comparative Survey [103]	2019	NL/LLM-Database	VLDB Journal	7.5	Optimize
S20	Consistency models in distributed systems: A survey on definiti... [120]	2019	Distributed DBs	arXiv	6	Operate, Optimize
S21	Frameworks for Querying Databases Using Natural Language: A Lit... [121]	2019	NL/LLM-Database	arXiv	6.5	Optimize
S22	Architecture, Security & Performance of DBaaS [77]	2019	Security/Privacy	Trans. Emerg. Telecom. Tech.	6.5	Design, Operate

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S23	A comparative analysis of adaptive consistency approaches in cl... [122]	2019	Distributed DBs	J. Parallel Distrib. Comput.	7.5	Operate, Optimize
S24	Multi-Model Databases: A New Journey [52]	2019	Multi-Model DBs	ACM Comp. Surv.	8	Design, Deploy
S25	LSM-based storage techniques: a survey [123]	2019	NoSQL/NewSQL	VLDB J.	7.5	Design, Optimize
S26	The Query Translation Landscape: a Survey [124]	2019	Multi-Model DBs	arXiv	6.5	Optimize
S27	Blockchains vs. Distributed Databases: Dichotomy and Fusion [125]	2019	Distributed DBs	arXiv	6	Design, Operate, Optimize
S28	NoSQL Databases: Yearning for Disambiguation [126]	2020	NoSQL/NewSQL	arXiv	6	Design, Operate
S29	Benchmarking Graph Data Management and Processing Systems: A Survey [127]	2020	Benchmarking	arXiv	6.5	Optimize
S30	SQL, NewSQL, and NoSQL Databases: A Comparative Survey [48]	2020	NoSQL/NewSQL	IEEE ICICS	6	Design, Deploy
S31	Real-time analytics, hybrid transactional/analytical processing... [128]	2020	HTAP Systems	Proc. ISP RAS	6.5	Design, Optimize
S32	Benchmarking learned indexes [129]	2020	AI4DB/Learned Opt.	PVLDB	7.5	Optimize
S33	Querying in the Age of Graph Databases and Knowledge Graphs [130]	2021	Graph Databases	ACM SIGMOD Record	6.5	Design, Optimize

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S34	Big Data Storage Tools Using NoSQL Databases and Their Applicat... [131]	2021	NoSQL/NewSQL	Computing and Informatics	8	Design, Operate, Optimize, Evolve
S35	Evaluation of Distributed Databases in Hybrid Clouds and Edge C... [132]	2021	Edge/Fog DBs	arXiv (Cornell University)	6.5	Design, Deploy, Optimize
S36	Comparison and evaluation of state-of-the-art LSM merge policies [133]	2021	NoSQL/NewSQL	VLDB J.	7.5	Optimize
S37	Cloud Computing Security: Threats & Mitigation—An SLR [75]	2021	Security/Privacy	IEEE Access	7	Operate
S38	Constructing and analyzing the LSM compaction design space [134]	2021	NoSQL/NewSQL	PVLDB	7.5	Design, Optimize
S39	The Forgotten Document-Oriented Database Management Systems: An... [135]	2021	NoSQL/NewSQL	Big Data Research	7.5	Design, Optimize
S40	Integrity Authentication for SQL Query Evaluation on Outsourced... [136]	2021	Security/Privacy	IEEE TKDE	7.5	Operate
S41	Towards Polyglot Data Stores – Overview and Open Research Ques... [137]	2022	Multi-Model DBs	arXiv	6.5	Design, Optimize
S42	From Data Warehouse to Lakehouse: A Comparative Review [10]	2022	Data Lakehouse	IEEE BigData	6.5	Design, Deploy
S43	A Survey of Comparison Different Cloud Database Performance: SQ... [138]	2022	NoSQL/NewSQL	Passer J.	7	Design, Deploy, Operate, Optimize

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S44	Database Meets AI: A Survey [9]	2022	AI4DB/Learned Opt.	IEEE TKDE	8	Optimize, Evolve
S45	Consistency issue and related trade-offs in distributed replica... [139]	2023	Distributed DBs	Radioelectron. Comput. Syst.	7	Design, Operate, Optimize
S46	Demystifying Graph Databases: Analysis and Taxonomy of Data Org... [140]	2023	Graph Databases	ACM Comput. Surv.	8	Design, Operate, Optimize
S47	Multidimensional Modelling in NoSQL Database : A Systematic Rev... [100]	2023	NoSQL/NewSQLCISTI		7.5	Design
S48	Multi-model query languages: taming the variety of big data [141]	2023	Multi-Model DBs	Distrib. Parallel Databases	7.5	Design, Optimize
S49	SQL and NoSQL Database Software Architecture Performance Analys... [142]	2023	NoSQL/NewSQLBig Data	Cogn. Comput.	8	Design, Deploy, Optimize
S50	A Survey on Property-Preserving Database Encryption Techniques ... [143]	2023	Security/Privacy	arXiv	6.5	Design, Operate, Optimize
S51	A Comprehensive Survey on Vector Database: Storage and Retrieval... [144]	2023	Vector Databases	arXiv	7.0	Design, Optimize, Operate
S52	The World of Graph Databases from An Industry Perspective [145]	2023	Graph Databases	ACM SIGMOD Rec.	7.5	Design, Operate, Optimize
S53	An Overview on Cloud Distributed Databases for Business Environ... [146]	2023	DBaaS	arXiv	6	Design, Deploy, Operate, Optimize

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S54	Learned Index: A Comprehensive Experimental Evaluation [94]	2023	AI4DB/Learned Opt.	PVLDB	7	Optimize
S55	Understanding Graph Databases: A Comprehensive Tutorial and Sur... [147]	2024	Graph Databases	arXiv	6	Design, Optimize
S56	Energy Consumption and Performance Evaluation of Multi-Model No... [148]	2024	Benchmarking	R. Inform. Teor. Apl.	7	Design, Optimize
S57	Retrieval-Augmented Generation for LLMs: A Survey [149]	2024	NL/LLM-Database	arXiv	7	Deploy, Operate
S58	Databases in Edge and Fog Environments: A Survey [13]	2024	Edge/Fog DBs	ACM Comp. Surv.	7.5	Design, Deploy, Operate
S59	HTAP Databases: A Survey [7]	2024	HTAP Systems	arXiv	6.5	Design, Operate, Optimize
S60	The evolution of data storage architectures: examining the secu... [150]	2024	Data Lakehouse	J. Data Inf. Manag.	7.5	Design, Operate, Evolve
S61	Cloud-Native Databases: A Survey [6]	2024	Cloud-Native Arch.	IEEE TKDE	8	Design, Deploy, Operate, Optimize
S62	A survey of LSM-Tree based Indexes, Data Systems and KV-stores [151]	2024	NoSQL/NewSQL	IEEE SCEECS	7.5	Design, Optimize
S63	A Systematic Review of Automated Classification for Simple and ... [152]	2024	AI4DB/Learned Opt.	Comput. Syst. Sci. Eng.	7.5	Optimize, Operate
S64	Survey of Vector Database Management Systems [5]	2024	Vector Databases	VLDB Journal	7.5	Design, Deploy, Optimize

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S65	Data Lakehouse for Time Series Data: An SLR [54]	2024	Data Lakehouse	IEEE BigData	6.5	Design, Operate
S66	Data Modeling for Connected Data – A systematic literature rev... [153]	2024	Graph Databases	arXiv	7	Design, Evolve
S67	Sharding Distributed Databases: A Critical Review [154]	2024	Distributed DBs	arXiv	6	Design, Operate, Optimize
S68	A survey on hybrid transactional and analytical processing [155]	2024	HTAP Systems	VLDB J.	7.5	Design, Optimize
S69	DBMS Performance Comparisons: An SLR [69]	2024	Benchmarking	J. Syst. Softw.	7.5	Optimize
S70	Vector database management systems: Fundamental concepts, use-c... [156]	2024	Vector Databases	Cogn. Syst. Res.	7.5	Design, Operate, Optimize
S71	Cassandra vs. MongoDB: A Systematic Review [87]	2024	NoSQL/NewSQL	IEEE BDAI	6	Design, Optimize
S72	Automatic Configuration Tuning on Cloud Database: A Survey [157]	2024	AI4DB/Learned Opt.	arXiv	6.5	Optimize, Operate
S73	A Survey of Learned Indexes for the Multi-dimensional Space [158]	2025	AI4DB/Learned Opt.	ACM Comput. Surv.	7.5	Design, Optimize
S74	Survey: On the Landscape of Graph Databases [159]	2025	Graph Databases	arXiv	7	Design, Deploy, Operate, Optimize
S75	Evaluating Performance and User Perceptions of Multi-Cloud Data... [71]	2025	DBaaS	IEEE ICAC	6.5	Deploy, Operate, Optimize

Continued on next page

ID	Title (Abbreviated)	Year	Primary Topic	Venue	QA	Lifecycle Phases
S76	When Large Language Models Meet Vector Databases: A Survey [160]	2025	NL/LLM-Database	IEEE AIxMM	7	Design, Optimize
S77	NoSQL Databases: A Survey [12]	2025	NoSQL/NewSQLARIMA		7.5	Design, Deploy
S78	Graph Neural Networks for Databases: A Survey [161]	2025	AI4DB/Learned Opt.	arXiv	6.5	Optimize
S79	Evaluating Learned Indexes in LSM-tree Systems: Benchmarks, Insi... [162]	2025	AI4DB/Learned Opt.	arXiv	6.5	Optimize
S80	How good are multi-dimensional learned indexes? An experimental... [163]	2025	Benchmarking	VLDB J.	7.5	Optimize
S81	Survey on Cloud Database Security: Cryptographic Techniques, In... [164]	2025	Security/Privacy	EPJ Web Conf.	7.5	Operate, Optimize
S82	NoSQL Data Warehouse Optimizing Models [99]	2025	NoSQL/NewSQLBig Data Research		6	Design, Optimize
S83	Learning database optimization techniques: the state-of-the-art... [165]	2025	AI4DB/Learned Opt.	Front. Comput. Sci.	7.5	Optimize
S84	Evaluating the Performance of NoSQL Databases for Big Data in C... [166]	2025	NoSQL/NewSQLHighTech J.	Innov.	7.5	Deploy, Optimize
S85	Toward Understanding Bugs in Vector Database Management Systems [167]	2025	Vector Databases	arXiv	7	Deploy, Operate, Optimize

References

- [1] Gartner, Inc., Magic quadrant for cloud database management systems, Gartner Research, accessed: 2025-01-15 (2024).
- [2] A. Verbitski, A. Gupta, D. Saha, M. Brahmadesam, K. Gupta, R. Mittal, S. Krishnamurthy, S. Maurice, T. Kharatishvili, X. Bao, Amazon Aurora: Design considerations for high throughput cloud-native relational databases, in: Proceedings of the ACM SIGMOD International Conference on Management of Data, 2017, pp. 1041–1052. doi:10.1145/3035918.3056101.
- [3] B. Dageville, T. Cruanes, M. Zukowski, V. Antonov, A. Avanes, J. Bock, J. Claybaugh, D. Engovatov, M. Hentschel, J. Huang, A. W. Lee, A. Motivala, A. Q. Munir, S. Pelley, P. Povinec, G. Rahn, S. Triantafyllis, P. Unterbrunner, The Snowflake elastic data warehouse, in: Proceedings of the ACM SIGMOD International Conference on Management of Data, 2016, pp. 215–226. doi:10.1145/2882903.2903741.
- [4] P. Antonopoulos, A. Budovski, C. Diaconu, A. H. Saenz, J. Hu, H. Kodavalla, D. Kossmann, S. Lingam, U. F. Minhas, N. Prakash, V. Purber, H. Qu, C. S. Ravella, K. Reisteter, S. Shrotri, D. Tang, V. Wakade, Socrates: The new SQL server in the cloud, in: Proceedings of the ACM SIGMOD International Conference on Management of Data, 2019, pp. 1743–1756. doi:10.1145/3299869.3314047.
- [5] J. J. Pan, J. Wang, G. Li, Survey of vector database management systems, The VLDB Journal 33 (5) (2024) 1371–1400. doi:10.1007/s00778-024-00864-x.
- [6] H. Dong, C. Zhang, G. Li, H. Zhang, Cloud-native databases: A survey, IEEE Transactions on Knowledge and Data Engineering 36 (12) (2024) 7772–7791. doi:10.1109/TKDE.2024.3397508.
- [7] C. Zhang, G. Li, J. Zhang, X. Zhang, J. Feng, HTAP databases: A survey, arXiv preprint arXiv:2404.15670 (2024). arXiv:2404.15670.
- [8] R. Angles, M. Arenas, P. Barceló, A. Hogan, J. Reutter, D. Vrgoc, Foundations of modern query languages for graph databases, ACM Computing Surveys 50 (5) (2017) 1–40. doi:10.1145/3104031.
- [9] X. Zhou, C. Chai, G. Li, J. Sun, Database meets artificial intelligence: A survey, IEEE Transactions on Knowledge and Data Engineering 34 (3) (2022) 1096–1116. doi:10.1109/TKDE.2020.2994641.
- [10] A. A. Harby, F. Zulkernine, From data warehouse to lakehouse: A comparative review, in: 2022 IEEE International Conference on Big Data (Big Data), IEEE, 2022, pp. 389–395. doi:10.1109/BigData55660.2022.10020719.
- [11] A. Davoudian, L. Chen, M. Liu, A survey on NoSQL stores, ACM Computing Surveys 51 (2) (2018) 1–43. doi:10.1145/3158661.
- [12] D. Kpekpassi, D. Faye, NoSQL databases: A survey, Revue Africaine de Recherche en Informatique et Mathématiques Appliquées (ARIMA) (2025). doi:10.46298/arima.13970.
- [13] L. M. M. Ferreira, F. Coelho, J. Pereira, Databases in edge and fog environments: A survey, ACM Computing Surveys 56 (11) (2024) 1–40. doi:10.1145/3666001.
- [14] B. Kitchenham, S. Charters, Guidelines for performing systematic literature reviews in software engineering, Technical Report EBSE-2007-01, Keele University and Durham University (2007).
- [15] M. Armbrust, A. Ghodsi, R. Xin, M. Zaharia, Lakehouse: A new generation of open platforms that unify data warehousing and advanced analytics, in: Proceedings of CIDR, 2021.
- [16] R. Taft, I. Sharif, A. Matei, N. VanBenschoten, J. Lewis, T. Grieger, K. Niemi, A. Woods, A. Biber, R. Poss, P. Bardea, A. Ranade, B. Darnell, B. Gruneir, J. Jaffray, L. Zhang, P. Mattis, CockroachDB: The resilient geo-distributed SQL database, in: Proceedings of the ACM SIGMOD International Conference on Management of Data, 2020, pp. 1493–1509. doi:10.1145/3318464.3386134.
- [17] W. Cao, Y. Zhang, X. Yang, F. Li, S. Wang, Q. Hu, X. Cheng, Z. Chen, Z. Liu, J. Fang, B. Wang, Y. Wang, H. Sun, Z. Yang, Z. Cheng, S. Chen, J. Wu, W. Hu, J. Zhao, Y. Gao, S. Cai, Y. Zhang, J. Tong, PolarDB serverless: A cloud native database for disaggregated data centers, Proceedings of the ACM SIGMOD International Conference on Management of Data (2021) 2477–2489. doi:10.1145/3448016.3457560.
- [18] D. Huang, Q. Liu, Q. Cui, Z. Fang, X. Ma, F. Xu, L. Shen, L. Tang, Y. Zhou, M. Huang, W. Wei, C. Liu, J. Zhang, J. Li, X. Wu, L. Song, R. Sun, S. Yu, L. Zhao, N. Cameron, L. Pei, X. Tang, TiDB: A Raft-based HTAP database, in: Proceedings of the VLDB Endowment, Vol. 13, 2020, pp. 3072–3084. doi:10.14778/3415478.3415535.
- [19] Z. Yang, C. Yang, F. Han, M. Zhuang, B. Yang, Z. Lin, H. Xu, G. Li, L. Ran, S. Wang, Y. Li, R. Pan, L. Guo, OceanBase: A 707 million tpmC distributed relational database system, in: Proceedings of the VLDB Endowment, Vol. 15, 2022, pp. 3385–3397. doi:10.14778/3554821.3554830.
- [20] F. Färber, S. K. Cha, J. Primsch, C. Bornhövd, S. Siber, W. Lehner, SAP HANA database: Data management for modern business applications, in: SIGMOD Record, Vol. 40, 2012, pp. 45–51. doi:10.1145/2094114.2094126.
- [21] M. Elbattah, M. Roushdy, M. Aref, A.-B. M. Salem, Large-scale ontology storage and query using graph database-oriented approach: The case of Freebase, in: 2015 IEEE Seventh International Conference on Intelligent Computing and Information Systems (ICICIS), IEEE, 2015.
- [22] M. E. Coimbra, L. Svitakova, D. Vrgoč, A. P. Francisco, L. Veiga, Survey: Graph databases, arXiv preprint arXiv:2505.24758 (2025). arXiv:2505.24758.
- [23] G. Xiao, D. Calvanese, R. Kontchakov, D. Lembo, A. Poggi, R. Rosati, M. Zakharyashev, Ontology-based data access: A survey, Proceedings of the International Joint Conference on Artificial Intelligence (2018) 5511–5519.
- [24] A. Halevy, F. Korn, N. F. Noy, C. Olston, N. Polyzotis, S. Roy, S. E. Whang, Goods: Organizing Google’s datasets, in: Proceedings of the ACM SIGMOD International Conference on Management of Data, 2016, pp. 795–806. doi:10.1145/2882903.2903730.
- [25] V. Leis, M. Kuschewski, Towards cost-optimal query processing in the cloud, in: Proceedings of the VLDB Endowment, Vol. 14, 2021, pp. 1606–1612. doi:10.14778/3461535.3461549.
- [26] V. Mateljan, D. Čišić, D. Ogrizović, Cloud database-as-a-service (DaaS) – ROI, in: Proceedings of the 33rd International Convention MIPRO, IEEE, 2010.
- [27] I. Müller, R. Marroquín, G. Alonso, Lambda: Interactive data analytics on cold data using serverless cloud infrastructure, in: Proceedings of the ACM SIGMOD International Conference on Management of Data, 2020, pp. 115–130. doi:10.1145/3318464.3389758.
- [28] European Parliament and Council of the European Union, Regulation (EU) 2016/679 of the European Parliament and of the Council (General Data Protection Regulation), Official Journal of the European Union (2016).

- [29] E. Politou, E. Alepis, C. Patsakis, Forgetting personal data and revoking consent under the GDPR: Challenges and proposed solutions, *Journal of Cybersecurity* 4 (1) (2018) tty001. doi:10.1093/cybsec/tty001.
- [30] E. Ferrari, Database as a service: Challenges and solutions for privacy and security, in: 2009 IEEE Asia-Pacific Services Computing Conference (APSCC), IEEE, 2009, pp. 46–51. doi:10.1109/APSCC.2009.5394141.
- [31] J. Köhler, K. Jünemann, H. Hartenstein, Confidential database-as-a-service approaches: taxonomy and survey, *Journal of Cloud Computing* 4 (1) (2015). doi:10.1186/s13677-014-0025-1.
- [32] F. Chang, J. Dean, S. Ghemawat, W. C. Hsieh, D. A. Wallach, M. Burrows, T. Chandra, A. Fikes, R. E. Gruber, Bigtable: A distributed storage system for structured data, *ACM Transactions on Computer Systems* 26 (2) (2008) 1–26. doi:10.1145/1365815.1365816.
- [33] S. Ramanathan, S. Goel, S. Alagumalai, Comparison of cloud database: Amazon's SimpleDB and Google's Bigtable, in: 2011 International Conference on Recent Trends in Information Systems (ReTIS), 2011. doi:10.1109/ReTIS.2011.6146861.
- [34] T. Dory, B. Mejías, P. Van Roy, N.-L. Tran, Measuring elasticity for cloud databases, in: CLOUD COMPUTING 2011: The Second International Conference on Cloud Computing, GRIDs, and Virtualization, IARIA, 2011.
- [35] R. Cattell, Scalable SQL and NoSQL data stores, *ACM SIGMOD Record* 39 (4) (2011) 12–27. doi:10.1145/1978915.1978919.
- [36] J. Han, H. E. G. Le, J. Du, Survey on NoSQL database, in: 2011 6th International Conference on Pervasive Computing and Applications (ICPCA), IEEE, 2011, pp. 363–366.
- [37] G. C. Deka, A survey of cloud database systems, *IT Professional* 16 (2) (2014) 50–57. doi:10.1109/MITP.2013.1.
- [38] A. T. Litchfield, J. Althouse, A systematic review of cloud computing, big data and databases on the cloud, in: AMCIS, 2014. URL <https://aisel.aisnet.org/amcis2014/ServiceSystems/GeneralPresentations/1>
- [39] I. Arora, A. Gupta, Cloud databases: A paradigm shift in databases, *International Journal of Computer Science Issues* 9 (4) (2012).
- [40] J. R. Lourenço, B. Cabral, P. Carreiro, M. Vieira, J. Bernardino, Choosing the right nosql database for the job: a quality attribute evaluation, *Journal of Big Data* 2 (1) (2015). doi:10.1186/s40537-015-0025-0.
- [41] C. Curino, E. P. Jones, R. A. Popa, N. Malviya, E. Wu, S. Madden, H. Balakrishnan, N. Zeldovich, Relational cloud: A database-as-a-service for the cloud, in: Proceedings of CIDR, 2011.
- [42] W. Lehner, K.-U. Sattler, Database as a service (DBaaS), in: 2010 IEEE 26th International Conference on Data Engineering (ICDE), IEEE, 2010, pp. 1216–1217.
- [43] W. A. Shehri, Cloud database database as a service, *International Journal of Database Management Systems (IJDMs)* 5 (2) (2013) 1–12. doi:10.5121/ijdms.2013.5201.
- [44] K. Grolinger, W. A. Higashino, A. Tiwari, M. A. M. Capretz, Data management in cloud environments: NoSQL and NewSQL data stores, *Journal of Cloud Computing: Advances, Systems and Applications* 2 (1) (2013) 22. doi:10.1186/2192-113X-2-22.
- [45] A. Pavlo, G. Angulo, J. Arulraj, H. Lin, J. Lin, L. Ma, P. Menon, T. C. Mowry, M. Perron, I. Quah, S. Santurkar, A. Tomasic, S. Toor, D. V. Aken, Z. Wang, Y. Wu, R. Xian, T. Zhang, Self-driving database management systems, in: Proceedings of CIDR, 2017.
- [46] J. C. Corbett, J. Dean, M. Epstein, A. Fikes, C. Frost, J. J. Furman, S. Ghemawat, A. Gubarev, C. Heiser, P. Hochschild, W. Hsieh, S. Kanthak, E. Kogan, H. Li, A. Lloyd, S. Melnik, D. Mwaure, D. Nagle, S. Quinlan, R. Rao, L. Rolig, Y. Saito, M. Szymaniak, C. Taylor, R. Wang, D. Woodford, Spanner: Google's globally distributed database, *ACM Transactions on Computer Systems* 31 (3) (2013) 1–22. doi:10.1145/2491245.
- [47] F. Gessert, W. Wingerath, S. Friedrich, N. Ritter, NoSQL database systems: a survey and decision guidance, *Computer Science – Research and Development* 32 (3–4) (2017) 353–365. doi:10.1007/s00450-016-0334-3.
- [48] T. N. Khasawneh, M. H. AL-Sahlee, A. A. Safia, SQL, NewSQL, and NoSQL databases: A comparative survey, in: 2020 11th International Conference on Information and Communication Systems (ICICS), IEEE, 2020, pp. 013–021. doi:10.1109/ICICS49469.2020.239513.
- [49] B. G. Tudorica, C. Bucur, A comparison between several NoSQL databases with comments and notes, in: 2011 RoEduNet International Conference 10th Edition: Networking in Education and Research, IEEE, 2011.
- [50] A. Bader, O. Kopp, M. Falkenthal, Survey and comparison of open source time series databases, in: Datenbanksysteme für Business, Technologie und Web (BTW 2017) – Workshopband, Lecture Notes in Informatics (LNI), Gesellschaft für Informatik, 2017. URL <https://dl.gi.de/handle/20.500.12116/922>
- [51] G. Özsoyoğlu, R. T. Snodgrass, Temporal and real-time databases: A survey, *IEEE Transactions on Knowledge and Data Engineering* 7 (4) (1995) 513–532. doi:10.1109/69.404027.
- [52] J. Lu, I. Holubová, Multi-model databases: A new journey to handle the variety of data, *ACM Computing Surveys* 52 (3) (2019) 1–38. doi:10.1145/3323214.
- [53] S. Melnik, A. Gubarev, J. J. Long, G. Romer, S. Shivakumar, M. Tolton, T. Vassilakis, H. Ahmadi, D. Delorey, S. Min, M. Pasumansky, J. Shute, Dremel: A decade of interactive SQL analysis at web scale, in: Proceedings of the VLDB Endowment, Vol. 13, 2020, pp. 3461–3472. doi:10.14778/3415478.3415568.
- [54] M. Pohl, D. Staegemann, K. Turowski, Data lakehouse for time series data: A systematic literature review, in: 2024 IEEE International Conference on Big Data (BigData), IEEE, 2024. doi:10.1109/BigData62323.2024.10825961.
- [55] S. Salehi, A. Selamat, H. Fujita, Systematic mapping study on granular computing, *Knowledge-Based Systems* 80 (2015) 78–97. doi:10.1016/j.knosys.2015.02.018.
- [56] G. Murtaza, L. Shuib, A. W. A. Wahab, G. Mujtaba, H. F. Nweke, M. A. Al-garadi, F. Zulfiqar, G. Raza, N. A. Azmi, Deep learning-based breast cancer classification through medical imaging modalities: State of the art and research challenges, *Artificial Intelligence Review* 53 (3) (2020) 1655–1720. doi:10.1007/s10462-019-09716-7.
- [57] A. Qazi, H. Fayaz, A. Wadi, R. G. Raj, N. Rahim, W. A. Khan, The artificial neural network for solar radiation prediction and designing solar systems: A systematic literature review, *Journal of Cleaner Production* 104 (2015) 1–12. doi:10.1016/j.jclepro.2015.04.041.
- [58] W. Saeed, C. Omlin, Explainable AI (XAI): A systematic meta-survey of current challenges and future opportunities, *Knowledge-Based Systems* 263 (2023) 110273. doi:10.1016/j.knosys.2023.110273.

- [59] M. J. Page, J. E. McKenzie, P. M. Bossuyt, I. Boutron, T. C. Hoffmann, C. D. Mulrow, et al., The PRISMA 2020 statement: an updated guideline for reporting systematic reviews, *BMJ* 372 (2021) n71. doi:10.1136/bmj.n71.
- [60] R. J. Thara, S. Shine, C. Sanu, Optimizing the performance of database as a service (DaaS) model – a distributed approach, in: 2013 Fourth International Conference on Computing, Communications and Networking Technologies (ICCCNT), IEEE, 2013.
- [61] L. Ma, D. V. Aken, A. Hefny, G. Mezerhane, A. Pavlo, G. J. Gordon, Query-based workload forecasting for self-driving database management systems, *Proceedings of the ACM SIGMOD International Conference on Management of Data* (2018) 631–645 doi:10.1145/3183713.3196908.
- [62] E. Brewer, CAP twelve years later: How the “rules” have changed, *Computer* 45 (2) (2012) 23–29. doi:10.1109/MC.2012.37.
- [63] D. J. Abadi, Consistency tradeoffs in modern distributed database system design: CAP is only part of the story, *Computer* 45 (2) (2012) 37–42. doi:10.1109/MC.2012.33.
- [64] P. Viotti, M. Vukolić, Consistency in non-transactional distributed storage systems, *ACM Computing Surveys* 49 (1) (2016) 1–34. doi:10.1145/2926965.
- [65] T. Rabl, M. Danisch, M. Frank, S. Schindler, H.-A. Jacobsen, Just can't get enough—synthesizing big data benchmarks, *IEEE BigData Congress* (2019).
- [66] R. Marcus, P. Negi, H. Mao, N. Tatbul, M. Alizadeh, T. Kraska, Bao: Making learned query optimization practical, *Proceedings of the ACM SIGMOD International Conference on Management of Data* (2021) 1275–1288 doi:10.1145/3448016.3452838.
- [67] T. Kraska, A. Beutel, E. H. Chi, J. Dean, N. Polyzotis, The case for learned index structures, in: *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 2018, pp. 489–504. doi:10.1145/3183713.3196909.
- [68] J. Zhang, Y. Liu, K. Zhou, G. Li, Z. Xiao, B. Cheng, J. Xing, Y. Wang, T. Cheng, L. Liu, M. Ran, Z. Li, An end-to-end automatic cloud database tuning system using deep reinforcement learning, *Proceedings of the ACM SIGMOD International Conference on Management of Data* (2019) 415–432 doi:10.1145/3299869.3300085.
- [69] T. Taipalus, Database management system performance comparisons: A systematic literature review, *Journal of Systems and Software* 208 (2024) 111872. doi:10.1016/j.jss.2023.111872.
- [70] M. T. Samarta, A. A. S. Gunawan, M. E. Syahputra, Systematic literature review and comparative performance analysis of SQL and NoSQL databases in big data applications, in: 2024 International Conference on Informatics, Multimedia, Cyber and Information System (ICIMCIS), IEEE, 2024. doi:10.1109/ICIMCIS63449.2024.10957463.
- [71] D. M. P. D. Daundasekara, et al., Evaluating performance and user perceptions of multi-cloud database-as-a-service solutions: A mixed-methods study, in: 2025 International Conference on Computing, 2025. doi:10.1109/ICAC69156.2025.11361452.
- [72] B. Furht, A. Escalante, *Handbook of cloud computing*, Springer (2010). doi:10.1007/978-1-4419-6524-0.
- [73] Z. S. Zubi, On distributed database security aspects, in: 2009 International Conference on Multimedia Computing and Systems (ICMCS), IEEE, 2009.
- [74] V. Waghmare, K. Gojre, A. Watpade, Approach to enhancing concurrent and self-reliant access to cloud database: A review, in: 2015 International Conference on Computational Intelligence and Communication Networks (CICN), IEEE, 2015. doi:10.1109/CICN.2015.158.
- [75] B. Alouffi, M. Hasnain, A. Alharbi, W. Alosaimi, H. Alyami, M. Ayaz, A systematic literature review on cloud computing security: Threats and mitigation strategies, *IEEE Access* 9 (2021) 57792–57807. doi:10.1109/ACCESS.2021.3073203.
- [76] K. Munir, Security model for cloud database as a service (DBaaS), in: 2015 International Conference on Cloud Technologies and Applications (CloudTech), IEEE, 2015.
- [77] F. A. Khan, M. Jamjoom, A. Ahmad, M. Asif, An analytic study of architecture, security, privacy, query processing, and performance evaluation of database-as-a-service, *Transactions on Emerging Telecommunications Technologies* 30 (12) (2019) e3814. doi:10.1002/ett.3814.
- [78] A. Acar, H. Aksu, A. S. Uluagac, M. Conti, A survey on homomorphic encryption schemes: Theory and implementation, *ACM Computing Surveys* 51 (4) (2018) 1–35. doi:10.1145/3214303.
- [79] R. R. Ravan, N. B. Idris, Z. Mehrabani, A survey on querying encrypted data for database as a service, in: 2013 International Conference on Cyber-Enabled Distributed Computing and Knowledge Discovery (CyberC), IEEE, 2013. doi:10.1109/CyberC.2013.12.
- [80] L. Ferretti, M. Colajanni, M. Marchetti, Distributed, concurrent, and independent access to encrypted cloud databases, *IEEE Transactions on Parallel and Distributed Systems* 25 (2) (2014) 437–446. doi:10.1109/TPDS.2013.154.
- [81] A. Cuzzocrea, C. Mastroianni, G. M. Grasso, Private databases on the cloud: Models, issues and research perspectives, in: 2016 IEEE International Conference on Big Data (Big Data), 2016, pp. 3656–3661. doi:10.1109/BigData.2016.7841032.
- [82] M. F. Gholami, F. Daneshgar, G. Low, G. Beydoun, Cloud migration process—a survey, evaluation framework, and open challenges, *Journal of Systems and Software* 120 (2016) 31–69. doi:10.1016/j.jss.2016.06.068.
- [83] V. Narasayya, S. Das, M. Syamala, B. Chandramouli, S. Chaudhuri, SQLVM: Performance isolation in multi-tenant relational database-as-a-service, *Proceedings of CIDR* (2015).
- [84] S. Jain, D. Moritz, B. Howe, High variety cloud databases, in: 2016 IEEE 32nd International Conference on Data Engineering Workshops (ICDEW), 2016, pp. 12–19. doi:10.1109/ICDEW.2016.7495609.
- [85] E. Baralis, A. D. Valle, P. Garza, C. Rossi, F. Scullino, SQL versus NoSQL databases for geospatial applications, in: 2017 IEEE International Conference on Big Data (Big Data), IEEE, 2017, pp. 3388–3397.
- [86] K. Mahmood, K. Orsborn, T. Risch, Comparison of NoSQL datastores for large scale data stream log analytics, in: 2019 IEEE International Conference on Smart Computing (SMARTCOMP), IEEE, 2019. doi:10.1109/SMARTCOMP.2019.00093.
- [87] H. Tu, Cassandra vs. MongoDB: A systematic review of two NoSQL data stores in their industry uses, in: 2024 IEEE 7th International Conference on Big Data and Artificial Intelligence (BDIAI), IEEE, 2024. doi:10.1109/BDIAI62182.2024.10692676.
- [88] W. Hendricks, Review of NoSQL data stores: Using a reactive three-tier application for software developers to achieve a high availability application design architecture, in: 2019 International Conference (IEEE), IEEE, 2019.

- [89] N. Tripathi, NoSQL database education: A review of models, tools and teaching methods, *Journal of Systems and Software* 226 (2025) 112391. doi:10.1016/j.jss.2025.112391.
- [90] M. Sheng, S. Wang, Y. Zhang, K. Wang, J. Wang, Y. Luo, R. Hao, MQRLD: A multimodal data retrieval platform with query-aware feature representation and learned index based on data lake, *Information Processing & Management* 62 (5) (2025) 104101. doi:10.1016/j.ipm.2025.104101.
- [91] C. Pahl, P. Jamshidi, O. Zimmermann, Architectural principles for cloud software, *ACM Transactions on Internet Technology* 18 (2) (2018) 1–23. doi:10.1145/3104028.
- [92] T. A. M. Phan, J. K. Nurminen, M. Di Francesco, Cloud databases for internet-of-things data, in: 2014 IEEE International Conference on Internet of Things (iThings), 2014, pp. 117–124. doi:10.1109/iThings.2014.26.
- [93] X. Zhao, K.-Y. Lam, Density based learned spatial index for clustered data, *Information Systems* 135 (2026) 102606. doi:10.1016/j.is.2025.102606.
- [94] Z. Sun, X. Zhou, G. Li, Learned index: A comprehensive experimental evaluation, *Proceedings of the VLDB Endowment* 16 (8) (2023) 1992–2004. doi:10.14778/3594512.3594528.
- [95] J. Tan, T. Zhang, F. Li, J. Chen, Q. Zheng, P. Zhang, H. Qiao, Y. Shi, W. Cao, R. Zhang, iBTune: Individualized buffer tuning for large-scale cloud databases, in: *Proceedings of the VLDB Endowment*, Vol. 12, 2019, pp. 1221–1234. doi:10.14778/3339490.3339503.
- [96] B. Hilprecht, C. Binnig, U. Röhm, Learning a partitioning advisor for cloud databases, in: *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*, 2020, pp. 143–157. doi:10.1145/3318464.3389704.
- [97] E. Dritsas, M. Trigka, Machine learning in modern database systems: Techniques, architectures, and deployment challenges, *Engineering Applications of Artificial Intelligence* 170 (2026) 114225. doi:10.1016/j.engappai.2026.114225.
- [98] International Energy Agency, Data centres and data transmission networks, <https://www.iea.org/energy-system/buildings/data-centres-and-data-transmission-networks>, accessed: 2025-01-15 (2024).
- [99] M. Mouhibha, A. Mabrouk, Nosql data warehouse optimizing models: A comparative study of column-oriented approaches, *Big Data Research* 40 (2025) 100523. doi:10.1016/j.bdr.2025.100523.
- [100] L. M. Ferreira, et al., Multidimensional modelling in NoSQL database: A systematic review, in: 2023 18th Iberian Conference on Information Systems and Technologies (CISTI), IEEE, 2023. doi:10.23919/CISTI58278.2023.10211592.
- [101] M. Mouhibha, A. Mabrouk, An empirically-driven clustering framework for NoSQL data warehouse conversion: Optimizing column family design from relational big data using HDBSCAN, *Information and Software Technology* 195 (2026) 108117. doi:10.1016/j.infsof.2026.108117.
- [102] M. Ma, Z. Yin, S. Zhang, S. Wang, C. Zheng, X. Jiang, H. Hu, C. Luo, Y. Li, N. Qiu, F. Li, C. Chen, D. Pei, Diagnosing root causes of intermittent slow queries in cloud databases, in: *Proceedings of the VLDB Endowment*, Vol. 13, 2020, pp. 1176–1189. doi:10.14778/3389133.3389136.
- [103] K. Affolter, K. Stockinger, A. Bernstein, A comparative survey of recent natural language interfaces for databases, *The VLDB Journal* 28 (5) (2019) 793–819. doi:10.1007/s00778-019-00567-8.
- [104] M. Pourreza, D. Rafiei, DIN-SQL: Decomposed in-context learning of text-to-SQL with self-correction, *Advances in Neural Information Processing Systems* 36 (2023).
- [105] D. Gao, H. Wang, Y. Li, X. Sun, Y. Qian, B. Ding, J. Zhou, Text-to-SQL empowered by large language models: A benchmark evaluation, *Proceedings of the VLDB Endowment* 17 (5) (2024) 1132–1145. doi:10.14778/3641204.3641221.
- [106] B. Wang, C. Ren, J. Yang, X. Liang, J. Bai, L. Chai, Z. Yan, Q.-W. Zhang, D. Yin, X. Sun, Z. Li, MAC-SQL: A multi-agent collaborative framework for text-to-SQL, *arXiv preprint arXiv:2312.11242* (2024).
- [107] T. Yu, R. Zhang, K. Yang, M. Yasunaga, D. Wang, Z. Li, J. Ma, I. Li, Q. Yao, S. Roman, Z. Zhang, D. Radev, Spider: A large-scale human-labeled dataset for complex and cross-database semantic parsing and text-to-SQL task, in: *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018, pp. 3911–3921. doi:10.18653/v1/D18-1425.
- [108] X. Zhou, G. Li, Z. Sun, Z. Liu, W. Chen, J. Wu, J. Liu, R. Feng, G. Zeng, D-bot: Database diagnosis system using large language models, in: *Proceedings of the VLDB Endowment*, Vol. 17, 2024, pp. 2514–2527. doi:10.14778/3675034.3675043.
- [109] S. Xue, C. Jiang, W. Shi, F. Cheng, K. Chen, H. Yang, Z. Zhang, J. He, H. Zhang, G. Wei, W. Zhao, F. Zhou, D. Qi, H. Yi, S. Liu, F. Chen, DB-GPT: Empowering database interactions with private large language models, *arXiv preprint arXiv:2312.17449* (2024).
- [110] J. Lao, Y. Wang, Y. Li, J. Wang, Y. Zhang, Z. Cheng, W. Chen, M. Tang, J. Wang, GPTuner: A manual-reading database tuning system via GPT-guided bayesian optimization, *Proceedings of the VLDB Endowment* 17 (8) (2024) 1939–1952. doi:10.14778/3659437.3659449.
- [111] J. Lee, W.-S. Han, R. Kasperovics, J.-H. Lee, An in-depth comparison of subgraph isomorphism algorithms in graph databases, *Proceedings of the VLDB Endowment* 6 (2) (2012) 133–144. doi:10.14778/2535568.2448946.
- [112] M. Abouzeq, A. Idrissi, Database-as-a-service for big data: An overview, *International Journal of Advanced Computer Science and Applications* 7 (1) (2016). doi:10.14569/ijacsa.2016.070124.
- [113] R. Guralnik, Some problems on graph databases, *Proceedings of the Institute for System Programming of the RAS* 28 (4) (2016) 193–216. doi:10.15514/ispras-2016-28(4)-12.
- [114] S. S. Moghadam, J. Darmont, G. Gavin, Enforcing privacy in cloud databases, *arXiv:1708.09171* (2017). doi:10.48550/arxiv.1708.09171.
- [115] A. Siddiqua, A. Karim, A. Gani, Big data storage technologies: a survey, *Frontiers of Information Technology & Electronic Engineering* 18 (8) (2017) 1040–1070. doi:10.1631/fitee.1500441.
- [116] Z. Ali, H. M. Kiran, W. Shahzad, Evolutionary algorithms for query op-timization in distributed database sys-tems: A review, *ADCAIJ: Advances in Distributed Computing and Artificial Intelligence Journal* 7 (3) (2018) 115–128. doi:10.14201/adcaij201873115128.
- [117] B. IANCU, T. M. GEORGESCU, Saving large semantic data in cloud: A survey of the main dbaas solutions, *Informatica Economica* 22 (1/2018) (2018) 5–16. doi:10.12948/issn14531305/22.1.2018.01.

- [118] P. N. S., K. P. K. N. P., N. P. K. M., A survey on graph database management techniques for huge unstructured data, *International Journal of Electrical and Computer Engineering (IJECE)* 8 (2) (2018) 1140. doi:10.11591/ijece.v8i2.pp1140-1149.
- [119] V. D. S. Segalin, C. F. Dorneles, M. A. R. Dantas, Dbas multitenancy, auto-tuning and sla maintenance in cloud environments: a brief survey, *iSys - Brazilian Journal of Information Systems* 11 (2) (2018) 30–42. doi:10.5753/isys.2018.362.
- [120] H. N. S. Aldin, H. Deldari, M. H. Moattar, M. R. Ghods, Consistency models in distributed systems: A survey on definitions, disciplines, challenges and applications, *arXiv*: (2019). doi:10.48550/arxiv.1902.03305.
- [121] H. S. Dar, M. I. Lali, M. U. Din, K. M. Malik, S. A. C. Bukhari, Frameworks for querying databases using natural language: A literature review, *arXiv:1909.01822* (2019). doi:10.48550/arxiv.1909.01822.
- [122] A. Khelaifa, S. Benharzallah, L. Kahloul, R. Euler, A. Laouid, A. Bounceur, A comparative analysis of adaptive consistency approaches in cloud storage, *Journal of Parallel and Distributed Computing* 129 (2019) 36–49. doi:10.1016/j.jpdc.2019.03.006.
- [123] C. Luo, M. J. Carey, Lsm-based storage techniques: a survey, *The VLDB Journal* 29 (1) (2019) 393–418. doi:10.1007/s00778-019-00555-y.
- [124] M. N. Mami, D. Graux, H. Thakkar, S. Scerri, S. Auer, J. Lehmann, The query translation landscape: a survey, *arXiv:1910.03118* (2019). doi:10.48550/arxiv.1910.03118.
- [125] P. Ruan, T. T. A. Dinh, D. Loghin, M. Zhang, G. Chen, Q. Lin, B. C. Ooi, Blockchains vs. distributed databases: Dichotomy and fusion, *arXiv*: (2019). doi:10.48550/arxiv.1910.01310.
- [126] C. Asaad, K. Baïna, M. Ghogho, Nosql databases: Yearning for disambiguation, *arXiv:2003.04074* (2020). doi:10.48550/arxiv.2003.04074.
- [127] M. Dayarathna, T. Suzumura, Benchmarking graph data management and processing systems: A survey, *arXiv:2005.12873* (2020). doi:10.48550/arxiv.2005.12873.
- [128] S. D. Kuznetsov, P. E. Velikhov, Q. Fu, Real-time analytics, hybrid transactional/analytical processing, in-memory data management, and non-volatile memory, 2020 Ivannikov Ispras Open Conference (ISPRAS) (2020) 78–90doi:10.1109/ispras51486.2020.00019.
- [129] R. Marcus, A. Kipf, A. van Renen, M. Stoian, S. Misra, A. Kemper, T. Neumann, T. Kraska, Benchmarking learned indexes, *Proceedings of the VLDB Endowment* 14 (1) (2020) 1–13. doi:10.14778/3421424.3421425.
- [130] M. Arenas, C. Gutierrez, J. F. Sequeda, Querying in the age of graph databases and knowledge graphs, *Proceedings of the 2021 International Conference on Management of Data* (2021) 2821–2828doi:10.1145/3448016.3457545.
- [131] A. Faridoon, M. Imran, Big data storage tools using nosql databases and their applications in various domains: A systematic review, *Computing and Informatics* 40 (3) (2021) 489–521. doi:10.31577/cai_2021_3_489.
- [132] Y. Mansouri, V. Prokhorenko, F. Ullah, M. A. Babar, Evaluation of distributed databases in hybrid clouds and edge computing: Energy, bandwidth, and storage consumption, *arXiv*: (2021). doi:10.48550/arxiv.2109.07260.
- [133] Q. Mao, S. Jacobs, W. Amjad, V. Hristidis, V. J. Tsotras, N. E. Young, Comparison and evaluation of state-of-the-art lsm merge policies, *The VLDB Journal* 30 (3) (2021) 361–378. doi:10.1007/s00778-020-00638-1.
- [134] S. Sarkar, D. Staratzis, Z. Zhu, M. Athanassoulis, Constructing and analyzing the lsm compaction design space, in: *Proceedings of the VLDB Endowment*, Vol. 14, 2021, pp. 2216–2229. doi:10.14778/3476249.3476274.
- [135] C.-O. Truică, E.-S. Apostol, J. Darmont, T. B. Pedersen, The forgotten document-oriented database management systems: An overview and benchmark of native xml dodbmses in comparison with json dodbmses, *Big Data Research* 25 (2021) 100205. doi:10.1016/j.bdr.2021.100205.
- [136] B. Zhang, B. Dong, W. H. Wang, Integrity authentication for sql query evaluation on outsourced databases: A survey, *IEEE Transactions on Knowledge and Data Engineering* 33 (4) (2021) 1601–1618. doi:10.1109/tkde.2019.2947061.
- [137] D. Glake, F. Kiehn, M. Schmidt, F. Panse, N. Ritter, Towards polyglot data stores – overview and open research questions, *arXiv:2204.05779* (2022). doi:10.48550/arxiv.2204.05779.
- [138] T. Shareef, K. Sharif, B. Rashid, A survey of comparison different cloud database performance: Sql and nosql, *Passer Journal of Basic and Applied Sciences* 4 (1) (2022) 45–57. doi:10.24271/psr.2022.301247.1104.
- [139] J. Ahmed, A. Karpenko, O. Tarasyuk, A. Gorbenco, A. Sheikh-Akbari, Consistency issue and related trade-offs in distributed replicated systems and databases: a review, *Radioelectronic and Computer Systems* (2) (2023) 171–179. doi:10.32620/reks.2023.2.14.
- [140] M. Besta, R. Gerstenberger, E. Peter, M. Fischer, M. Podstawski, C. Barthels, G. Alonso, T. Hoefer, Demystifying graph databases: Analysis and taxonomy of data organization, system designs, and graph queries, *ACM Computing Surveys* 56 (2) (2023) 1–40. doi:10.1145/3604932.
- [141] Q. Guo, C. Zhang, S. Zhang, J. Lu, Multi-model query languages: taming the variety of big data, *Distributed and Parallel Databases* 42 (1) (2023) 31–71. doi:10.1007/s10619-023-07433-1.
- [142] W. Khan, T. Kumar, C. Zhang, K. Raj, A. M. Roy, B. Luo, Sql and nosql database software architecture performance analysis and assessments—a systematic literature review, *Big Data and Cognitive Computing* 7 (2) (2023) 97. doi:10.3390/bdcc7020097.
- [143] J. Koppenwallner, E. Schikuta, A survey on property-preserving database encryption techniques in the cloud, *arXiv:2312.12075* (2023). doi:10.48550/arxiv.2312.12075.
- [144] L. Ma, R. Zhang, Y. Han, S. Yu, Z. Wang, Z. Ning, J. Zhang, P. Xu, P. Li, Z. Qiao, W. Ju, C. Chen, D. Wang, K. Liu, P. Wang, P. Wang, Y. Fu, C. Liu, Y. Zhou, C.-T. Lu, A comprehensive survey on vector database: Storage and retrieval technique, challenge, *arXiv:2310.11703* (2023). doi:10.48550/arxiv.2310.11703.
- [145] Y. Tian, The world of graph databases from an industry perspective, *ACM SIGMOD Record* 51 (4) (2023) 60–67. doi:10.1145/3582302.3582320.
- [146] A. Vikiru, M. Muiruri, I. Ateya, An overview on cloud distributed databases for business environments, *arXiv:2301.10673* (2023).
- [147] S. Anuyah, V. Bolade, O. Agbaakin, Understanding graph databases: A comprehensive tutorial and survey, *arXiv:2411.09999* (2024). doi:10.48550/arxiv.2411.09999.

- [148] F. Falcão, J. Moura, G. Silva, C. Araujo, E. Sousa, E. Tavares, Energy consumption and performance evaluation of multi-model nosql dbms, *Revista de Informática Teórica e Aplicada* 30 (2) (2024) 132–140. doi:10.22456/2175-2745.136568.
- [149] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, M. Wang, H. Wang, Retrieval-augmented generation for large language models: A survey, *arXiv preprint arXiv:2312.10997* (2024).
- [150] N. Janssen, T. Ilayperuma, J. Jayasinghe, F. Bukhsh, M. Daneva, The evolution of data storage architectures: examining the secure value of the data lakehouse, *Journal of Data, Information and Management* 6 (4) (2024) 309–334. doi:10.1007/s42488-024-00132-1.
- [151] S. Mishra, A survey of lsm-tree based indexes, data systems and kv-stores, in: *2024 IEEE International Students' Conference on Electrical, Electronics and Computer Science (SCEecs)*, 2024, pp. 1–6. doi:10.1109/sceecs61402.2024.10482249.
- [152] R. A. Kadir, E. S. M. Surin, M. R. Sarker, A systematic review of automated classification for simple and complex query sql on nosql database, *Computer Systems Science and Engineering* 48 (6) (2024) 1405–1435. doi:10.32604/csse.2024.051851.
- [153] V. Santos, Data modeling for connected data – a systematic literature review, *arXiv:2410.10081* (2024). doi:10.48550/arxiv.2410.10081.
- [154] S. Solat, Sharding distributed databases: A critical review, *arXiv*: (2024). doi:10.48550/arxiv.2404.04384.
- [155] H. Song, W. Zhou, H. Cui, X. Peng, F. Li, A survey on hybrid transactional and analytical processing, *The VLDB Journal* 33 (5) (2024) 1485–1515. doi:10.1007/s00778-024-00858-9.
- [156] T. Taipalus, Vector database management systems: Fundamental concepts, use-cases, and current challenges, *Cognitive Systems Research* 85 (2024) 101216. doi:10.1016/j.cogsys.2024.101216.
- [157] L. Zhang, M. A. Babar, Automatic configuration tuning on cloud database: A survey, *arXiv:2404.06043* (2024). doi:10.48550/arxiv.2404.06043.
- [158] A. Al-Mamun, H. Wu, Q. He, J. Wang, W. G. Aref, A survey of learned indexes for the multi-dimensional space, *ACM Computing Surveys* 58 (4) (2025) 1–37. doi:10.1145/3768575.
- [159] M. E. Coimbra, L. Svitáková, D. Vrgoč, A. P. Francisco, L. Veiga, Survey: On the landscape of graph databases, *arXiv:2505.24758* (2025). doi:10.48550/arxiv.2505.24758.
- [160] Z. Jing, Y. Su, Y. Han, When large language models meet vector databases: A survey, in: *2025 Conference on Artificial Intelligence x Multimedia (AIxMM)*, 2025, pp. 7–13, *arXiv:2402.01763*. doi:10.1109/AIxMM62960.2025.00008.
- [161] Z. Li, Y. Li, Y. Luo, G. Li, C. Zhang, Graph neural networks for databases: A survey, *arXiv:2502.12908* (2025). doi:10.48550/arxiv.2502.12908.
- [162] J. Liu, J. Ye, M. Chen, M. Li, S. Luo, Evaluating learned indexes in lsm-tree systems: Benchmarks, insights and design choices, *arXiv*: (2025). doi:10.48550/arxiv.2506.08671.
- [163] Q. Liu, M. Li, Y. Zeng, Y. Shen, L. Chen, How good are multi-dimensional learned indexes? an experimental survey, *The VLDB Journal* 34 (2) (2025). doi:10.1007/s00778-024-00893-6.
- [164] R. V. Mahule, S. M. Jadhav, Survey on cloud database security: Cryptographic techniques, intrusion detection and intelligent defense mechanisms, in: *EPJ Web Conf.*, 2025. doi:10.1051/epjconf/202534101015.
- [165] S.-J. Qiao, H.-L. Fan, N. Han, L. Du, Y.-H. Peng, R.-M. Tang, X. Qin, Learning database optimization techniques: the state-of-the-art and prospects, *Frontiers of Computer Science* 19 (12) (2025). doi:10.1007/s11704-025-41116-7.
- [166] A. E. Topcu, A. M. Rmis, Y. I. Alzoubi, A. O. Çibikdiken, D. Chowdhury, E. Elbasi, M. Abdel-Maguid, S. Almadhun, Evaluating the performance of nosql databases for big data in cloud computing environments, *HighTech and Innovation Journal* 6 (3) (2025) 808–830. doi:10.28991/hij-2025-06-03-05.
- [167] Y. Xie, X. Hou, Y. Zhao, S. Wang, K. Chen, H. Wang, Toward understanding bugs in vector database management systems, *arXiv*: (2025). doi:10.48550/arxiv.2506.02617.